

3_Hazim_Pima_Indians_Diabetes_93_4_Acc,_90_3_val_Acc

February 10, 2021

1 The Pima Indians Diabetes Dataset

2 Step 1: Answering the question

2.1 Loading the Dataset

```
[ ]: # Get the dataset by raw URL
!wget https://github.com/hazmash5/ds-projects/raw/main/Proj_05_Diabetes/data/
    ↪pima-indians-diabetes.csv
```

```
--2021-02-08 12:14:03-- https://github.com/hazmash5/ds-
projects/raw/main/Proj_05_Diabetes/data/pima-indians-diabetes.csv
Resolving github.com (github.com)... 140.82.113.4
Connecting to github.com (github.com)|140.82.113.4|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://raw.githubusercontent.com/hazmash5/ds-
projects/main/Proj_05_Diabetes/data/pima-indians-diabetes.csv [following]
--2021-02-08 12:14:03-- https://raw.githubusercontent.com/hazmash5/ds-
projects/main/Proj_05_Diabetes/data/pima-indians-diabetes.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
151.101.0.133, 151.101.64.133, 151.101.128.133, ...
Connecting to raw.githubusercontent.com
(raw.githubusercontent.com)|151.101.0.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 23278 (23K) [text/plain]
Saving to: pima-indians-diabetes.csv

pima-indians-diabet 100%[=====>] 22.73K --.-KB/s in 0s

2021-02-08 12:14:04 (80.7 MB/s) - pima-indians-diabetes.csv saved
[23278/23278]
```

```
[ ]: # Create new folder and name it data
!mkdir data
```

```
[ ]: # Moving our datasets to the data folder
!mv pima-indians-diabetes.csv data/

[ ]: # Showing the first line of the dataset
!head -n 3 data/pima-indians-diabetes.csv
```

```
6,148,72,35,0,33.6,0.627,50,1
1,85,66,29,0,26.6,0.351,31,0
8,183,64,0,0,23.3,0.672,32,1
```

We can notes the following: 1. There is no header row so header= None must be used while we read the csv. 2. There is no need to use Sep parameter because the separation between the values is (,) as the default separation of Panda csv Sep.

```
[ ]: # Showing the number of lines.
!cat data/pima-indians-diabetes.csv | wc -l
```

```
767
```

We have 768 of instances

2.2 Introduction

The Pima Indians Diabetes Dataset involves predicting the onset of diabetes within five years in Pima Indians given medical details. It is a binary (2-class) classification problem. Several familiar types of classification models algorithms utilized: 1. To choose the best classification algorithms and efficiently perform another appropriate comparison between the same algorithms. 2. To compare the utilized feature engineering and pre-processing methods. 3. To get a broad range of choices.

Utilized classification models are respectively: 1. Logistic Regression 2. Linear Discriminant Analysis 3. K Neighbors Classifier 4. Decision Tree Classifier 5. Gaussian NB 6. Support Vector Classifier 7. XGBoost Classifier

2.3 Required libraries

This notebook uses several Python packages that come standard with the Google Colaboratory. The primary libraries that we'll be operating are respectively: * **NumPy**: Provides a fast numerical array structure and helper functions. * **Pandas**: Provides a DataFrame structure to store data in memory and work with it easily and efficiently. * **Scikit-learn**: The essential Machine Learning package in Python. * **XGBoost**: Optimized distributed gradient boosting library designed to be highly efficient, flexible and portable matplotlib: Basic plotting library in Python; most other Python plotting libraries are built on top of it. * **Seaborn**: Advanced statistical plotting library. * **watermark**: A Jupyter Notebook extension for printing timestamps, version numbers, and hardware information.

```
[ ]: import pandas as pd
import numpy as np

%matplotlib inline
```

```
import matplotlib.pyplot as plt
import seaborn as sb
```

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from xgboost import XGBClassifier
models = []
models.append(('LR', LogisticRegression()))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('SVM', SVC()))
models.append(('XGB', XGBClassifier()))
```

2.4 The problem domain

Our company just got funded to create a smartphone app that automatically female diabetes detection to use in remote villages in India, from simple devices for each test attribute and fill it in the smartphone, we will be building part of the data analysis pipeline for this app. We tasked by the Head of Data Science to create a machine learning model, the model takes eight attributes from the user and detects diabetes based on those attributes alone. We got a dataset from the field researchers to develop the model, which includes predicting the onset of diabetes within five years in Pima-Indians given medical details. With the following attributes: * Number of times pregnant. * Plasma glucose concentration a 2 hours in an oral glucose tolerance test. * Diastolic blood pressure (mm Hg). * Triceps skinfold thickness (mm). * 2-Hour serum insulin (mu U/ml). * Body mass index (weight kg / height m2). * Diabetes pedigree function. * Age (years). * Class variable (0 or 1).

2.5 Data analysis checklist:

The data analysis checklist: 1. **Specify the type of data analytic question (e.g. exploration, association causality) before touching the data:** We are trying to detect female diabetes tests (Positive test or Negative test) based on eight continuous attributes. 2. **Define the metric for success before beginning:** We will use the accuracy to quantify how well our model is performing. they told us that we should achieve at least 77% accuracy.

3 Step 2: EDA Exploratory Data Analysis

```
[ ]: df= pd.read_csv('/content/data/pima-indians-diabetes.csv', header= None)
df.columns= ['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age', 'class']
df
```

```
[ ]:      preg  plas  pres_mm  skin_mm  insu   BMI   DPF  age  class
0         6   148      72      35     0  33.6  0.627  50     1
1         1    85      66      29     0  26.6  0.351  31     0
2         8   183      64       0     0  23.3  0.672  32     1
3         1    89      66      23    94  28.1  0.167  21     0
4         0   137      40      35   168  43.1  2.288  33     1
..      ...   ...      ...      ...   ...   ...   ...   ...   ...
763      10   101      76      48   180  32.9  0.171  63     0
764       2   122      70      27     0  36.8  0.340  27     0
765       5   121      72      23   112  26.2  0.245  30     0
766       1   126      60       0     0  30.1  0.349  47     1
767       1    93      70      31     0  30.4  0.315  23     0
```

[768 rows x 9 columns]

```
[ ]: # Check duplicated rows
print('The number of repeated rows= ',
      df.duplicated(keep='last').sum() )
```

The number of repeated rows= 0

There is no duplicated rows

```
[ ]: print('The number of null values= ',
      df.isnull().sum().sum())
```

The number of null values= 0


There is no NaN values.

3.1 1. Using pie plot

By using a pie plot to visualize and compute the difference between the categories, we can perceive the notable difference between the numbers of categories. A balance issue must be considered.

```
[ ]: color= ['#9999ff', '#ff4d4d']
[ ]: pd.value_counts(df['class']).plot(kind='pie',
                                     autopct='%.2f', fontsize=15,
                                     colors= color,
                                     title='Categories percentage')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc14a1524a8>
```




By using pie plot to visualize and compute the difference between the categories; we can notes the big difference between numbers of categories. There is an unbalance issue must be considered.

3.2 2. Using box plot

By using a box plot to receive an indication of how the values in the data are spread out, and to visualize the distribution of values within each attribute, we can notes the following: 1. All attributes values spread between 0 and 200 except insulin values, 2. The variance of values is extremely, so we must utilize preprocessing methods before training.

```
[ ]: color = dict(boxes='DarkGreen', whiskers='DarkOrange', medians='DarkBlue',
    ↪ caps='Gray')
df[df.columns[:-1]].plot(kind='box', color=color, sym='r+', figsize=(9,3),
    ↪ rot=30)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc13f6d0898>
```




3.3 3. Using Andrews curves

By exploiting Andrew's curves plot to visualize data clustering for each class, we can notes: 1. Curves belonging to samples of a similar class aren't closer together. 2. The curves of the two classes mix together and don't define structures. 3. It is problematic to target those classes, add features must be considered.


```
[ ]: pd.plotting.andrews_curves(df, 'class', color=color)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc14a050e10>
```



```
[ ]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
all_inputs = sc.fit_transform(all_inputs)
ddf = pd.DataFrame(all_inputs)
ddf['class'] = all_labels
pd.plotting.andrews_curves(ddf, 'class', color=color)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc149f502e8>
```

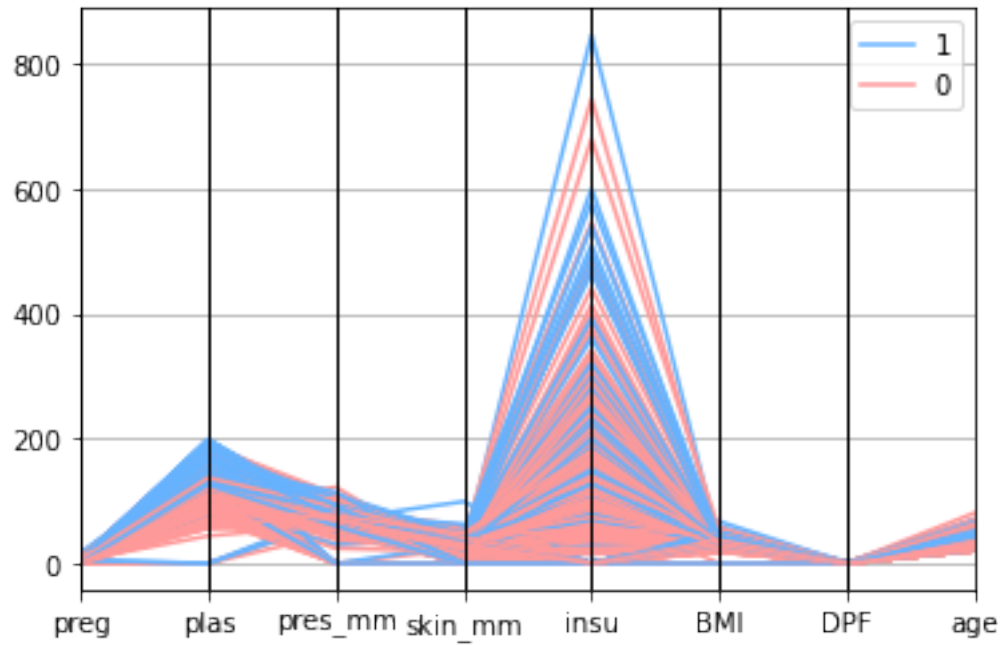


3.4 4. Using Parallel coordinates plot

By using Parallel coordinates plot to comparing variables together and observing the relationships between them, we can notes that there are no significant phenomena for each class, between the attributes.

```
[ ]: pd.plotting.parallel_coordinates(df, 'class', color=['#66b3ff', '#ff9999'])
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd10b6b4cc0>
```

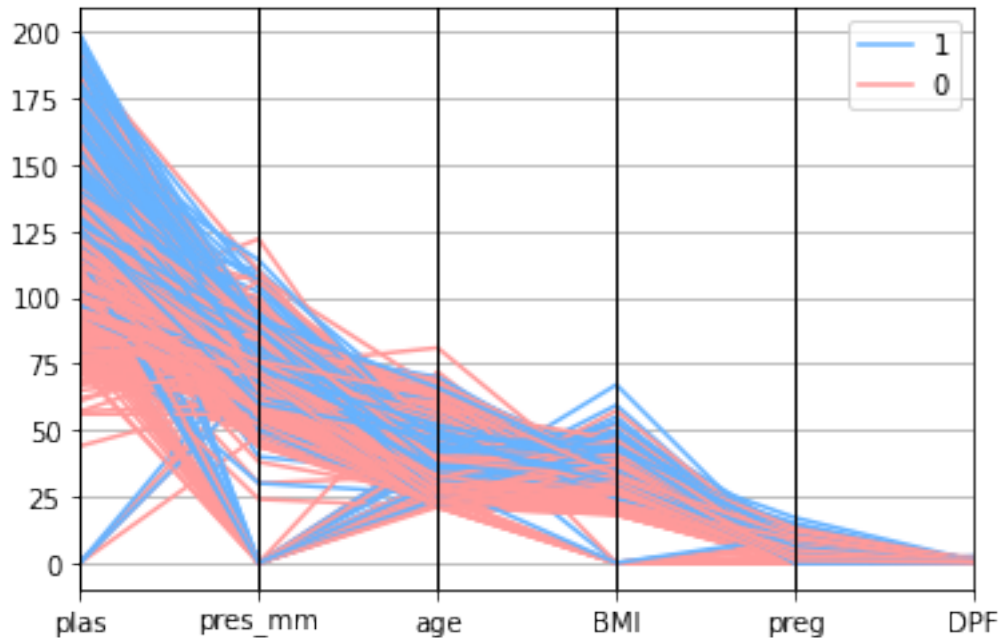


```
[ ]: df.columns
```

```
[ ]: Index(['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age',
          'class'],
          dtype='object')
```

```
[ ]: pd.plotting.parallel_coordinates(df[['plas', 'pres_mm', 'age', 'BMI', 'preg', 'DPF', 'class']], 'class', color=['#66b3ff', '#ff9999'])
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd0f41da4a8>
```





3.5 5. Using Radviz Plot

By using Radviz plotting to recognize clustering attribute for each class, we can notes the following: 1. Classes are clustering to the same attributes. 2. There are outlier instances, so outlier detection methods must be used. 3. Some attributes do not affect the categories, so feature selection methods must be used.

```
[ ]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler()
all_inputs = sc.fit_transform(all_inputs)
ddf = pd.DataFrame(all_inputs)
ddf.columns = df.columns[0:-1]
plt.figure(figsize=(9, 9))
ddf['class'] = all_labels
pd.plotting.radviz(ddf, 'class', color=['#ff9999', '#66b3ff'], )
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd10b6e2208>
```




There is no clear recognize distinguished cluster for each class to attributes as we saw in the previous figures. We can see a lot of outliers corresponding to the negative tests class.

3.6 6. Using hist plot

By using histogram plots we can visualize mean, median, standard deviation, and mode for the values of the attributes.


```
[ ]: df.describe().T['mean'].hist()
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f54a417ee48>
```




```
[ ]: df.median().hist()
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc13eed7cf8>
```



```
[ ]: df.describe().T['std'].hist()
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f54a417ecf8>
```




The mode is the most frequent observation

```
[ ]: fig = plt.figure(figsize = (10,10))
     ax = fig.gca()
     df.mode().hist(ax=ax)
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: UserWarning: To output multiple subplots, the figure containing the passed axes is being cleared
This is separate from the ipykernel package so we can avoid doing imports until

```
[ ]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f538ef0>,
           <matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f4b2630>,
           <matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f462898>],
          [<matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f494b00>,
           <matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f4dbc18>,
           <matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f3fd400>],
          [<matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f3b0668>,
           <matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f366898>,
           <matplotlib.axes._subplots.AxesSubplot object at 0x7fc13f366908>]],
          dtype=object)
```



3.7 7. Using Violin plot

By using violin plot to shows the distribution of quantitative data across several levels of categorical variables such that those distributions can be compared and features a kernel density estimation of the underlying distribution we can notes

1. Many zeros values.
2. The mean values for each class are different, so we must consider this when imputing Nan values.
3. There are outlier values.

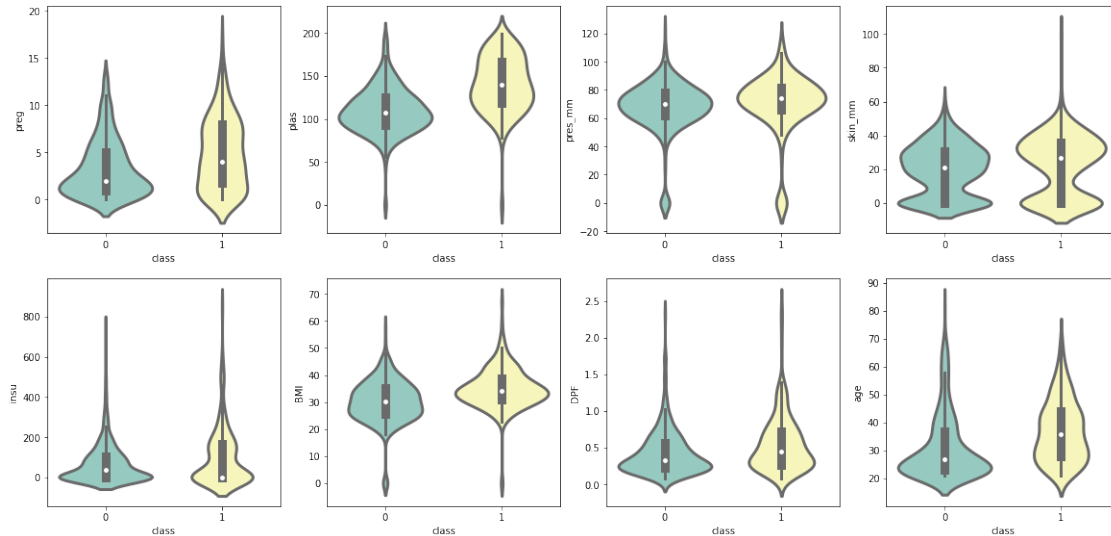
```
[ ]: plt.figure(figsize=(20, 20))

for column_index, column in enumerate(df):
```

```

if column == 'class':
    continue
plt.subplot(4, 4, column_index + 1)
sb.violinplot(x='class', y=column, data=df, inner="box", palette="Set3",
→cut=2, linewidth=3)

```



After talking with the field researchers, they fill the null values with zero so we must replace all not logical zeros values with NaN values.

3.8 8. Using Pair and KDE Kernel Density Estimate Plot

By using Pair and KDE plot to visualize distribution of single variables and relationships between variables we can notes 1. Relationships between some attribute, 2. A lot of zeros also, 3. Probability distributions are close and same.

```

[:]: # Function to calculate correlation coefficient between two arrays
def corr(x, y, **kwargs):

    # Calculate the value
    coef = np.corrcoef(x, y)[0][1]
    # Make the label
    label = r'$\rho$ = ' + str(round(coef, 2))

    # Add the label to the plot
    ax = plt.gca()
    ax.annotate(label, xy = (0.2, 0.95), size = 20, xycoords = ax.transAxes)

# Create a pair grid instance
grid = sb.PairGrid(data= df,
                    vars = ['No_preg', 'plas', 'pres_mm', 'skin_mm',


```

```

        'insu', 'BMI', 'DPF', 'age', 'class'], height = 5)

# Map the plots to the locations
grid = grid.map_upper(plt.scatter, color = 'red')
grid = grid.map_upper(corr)
grid = grid.map_lower(sb.kdeplot, cmap = 'Reds',)
grid = grid.map_diag(plt.hist, bins = 10, edgecolor = 'k', color = 'darkred');


```



3.9 9. Using the Correlation Matrix Heat map

By using the Correlation Matrix Heat map (figure 2.9) to illustrate the relationship between variables, we can note no significant case of multicollinearity is observed because all of the correlation coefficients are less than 0.7.

```
[ ]: plt.figure(figsize=(9, 7))
matrix = np.triu(df.corr())
corrMatrix = df.corr()
sb.heatmap(corrMatrix, annot=True, cmap='Blues', linewidths=0.15,
           mask=matrix, cbar_kws= {'orientation': 'horizontal'})
plt.show()
```



4 Step 3: Tidying the data

5 1. Fill NaN Methods Comparison

```
[ ]: df= pd.read_csv('/content/data/pima-indians-diabetes.csv', header= None)
df.columns= ['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age', 'class']
```

```
[ ]: df.columns
```

```
[ ]: Index(['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age',
          'class'],
          dtype='object')
```

```
[ ]: (df['age']==0).sum()
```

```
[ ]: 0
```



```
[ ]: print(''Columns could't be zero are: '' , list(df.columns[1:6]))
```

Columns could't be zero are: ['plas', 'pres_mm', 'skin_mm', 'insu', 'BMI']

```
[ ]: # Redefining zeros into NaN values
df[df.columns[1:6]] = df[df.columns[1:6]].replace(0, np.nan)
```

```
[ ]: df[df.columns[1:6]].isnull().sum()
```

```
[ ]: plas      5
     pres_mm   35
     skin_mm  227
     insu     374
     BMI      11
     dtype: int64
```

I will compare between 5 technique to fill NaN values

5.1 1. Remove Rows With Missing Values

```
[ ]: df= pd.read_csv('/content/data/pima-indians-diabetes.csv', header= None)
df.columns= ['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age', '
→'class']
df[df.columns[1:6]] = df[df.columns[1:6]].replace(0, np.nan)
```

```
[ ]: df.dropna(inplace=True)
df.shape
```

```
[ ]: (392, 9)
```

```
[ ]: print('Percentage of dropped data=', (768-392)/768)
```

Percentage of dropped data= 0.4895833333333333

```
[ ]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
print(all_inputs.shape)
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
                                                    test_size=0.25,
→random_state=1,)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_1 = []; names = []; seed=42
for name, model in models:
```

```

kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
cv_results_1 = cross_val_score(model, X_train, y_train, cv=kfold,
→scoring='accuracy' )
results_1.append(cv_results_1)
names.append(name)
print(f"{name}, {cv_results_1.mean()}, {cv_results_1.std()}))")

```

```

(392, 8)
LR, 0.7861904761904762, 0.12857054673418924))
LDA, 0.7897619047619047, 0.12802595931491162))
KNN, 0.7623809523809524, 0.112927444444886092))
CART, 0.6897619047619046, 0.1308464939868934))
NB, 0.7797619047619049, 0.11212111810322156))
SVM, 0.7554761904761904, 0.11571795996518176))
XGB, 0.7816666666666667, 0.11272018010217857))

```

5.2 2. Impute Missing Values (Mean Value Filling)

```

[:]: df= pd.read_csv('/content/data/pima-indians-diabetes.csv', header= None)
df.columns= ['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age',
→'class']
df[df.columns[1:6]]= df[df.columns[1:6]].replace(0, np.nan)

df1 = df.loc[df['class'] == 1]
df2 = df.loc[df['class'] == 0]

```

```

[:]: df1.fillna(df1.mean(), inplace=True)
df2.fillna(df2.mean(), inplace=True)

dataframe = [df1, df2]
df = pd.concat(dataframe)

```

```

[:]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
all_inputs.shape

```

```

[:]: (768, 8)

```

```

[:]: (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
test_size=0.25,
→random_state=1,
)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_2 = []; names = []; seed=42
for name, model in models:

```

```

kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
cv_results_2 = cross_val_score(model, X_train, y_train, cv=kfold,
→scoring='accuracy')
results_2.append(cv_results_2)
names.append(name)
print(f"{name}, {cv_results_2.mean()}, {cv_results_2.std()}))")

```

```

LR, 0.7738300492610837, 0.0689196220881071))
LDA, 0.7737684729064039, 0.07106786364748625))
KNN, 0.8173645320197045, 0.07886207573510806))
CART, 0.8796798029556652, 0.06851110409920169))
NB, 0.7720443349753694, 0.0700920699840484))
SVM, 0.8349137931034484, 0.06586019332396725))
XGB, 0.8869458128078817, 0.05642168499262387))

```

5.3 3. Impute Missing Values (Median Value Filling)

```

[:]: df= pd.read_csv('/content/data/pima-indians-diabetes.csv', header= None)
df.columns= ['No_preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF',
→'age', 'class']
df[df.columns[1:6]] = df[df.columns[1:6]].replace(0, np.nan)

```

```

df1 = df.loc[df['class'] == 1]
df2 = df.loc[df['class'] == 0]

```

```

[:]: df1.fillna(df1.median(), inplace=True)
df2.fillna(df2.median(), inplace=True)

```

```

dataframe = [df1, df2]
df = pd.concat(dataframe)

```

```

[:]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
all_inputs.shape

```

```

[:]: (768, 8)

```

```

[:]: (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
→test_size=0.25,

```

```

→random_state=1)

```

```

from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

```

```

results_3 = []; names = []; seed=42

```

```

for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)

```

```

cv_results_3 = cross_val_score(model, X_train, y_train, cv=kfold,
→scoring='accuracy')
results_3.append(cv_results_3)
names.append(name)
print(f"{name}, {cv_results_3.mean()}, {cv_results_3.std()}))")

```

```

LR, 0.7633004926108373, 0.07843541820685897))
LDA, 0.7703817733990148, 0.07992931047044209))
KNN, 0.8070812807881772, 0.07939031074825631))
CART, 0.8729064039408868, 0.048698449505117515))
NB, 0.7598522167487685, 0.08545857474403694))
SVM, 0.8504926108374384, 0.05873674193644676))
XGB, 0.888793103448276, 0.06942809765514414))

```

5.4 4. Using back Filling

```

[: df= pd.read_csv('/content/data/pima-indians-diabetes.csv', header= None)
df.columns= ['No_preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF',
→'age', 'class']
df[df.columns[1:6]] = df[df.columns[1:6]].replace(0, np.nan)

```

```

df1 = df.loc[df['class'] == 1]
df2 = df.loc[df['class'] == 0]

```

```

[: df1.fillna(method= 'backfill', inplace=True)
df2.fillna(method= 'backfill', inplace=True)

```

```

df1.fillna(df1.mean(), inplace=True)
df2.fillna(df2.mean(), inplace=True)

```

```

dataframe = [df1, df2]
df = pd.concat(dataframe)

```

```

[: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
all_inputs.shape

```

```

[: (768, 8)

```

```

[: (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
test_size=0.25,

```

```

→random_state=1)

```

```

# Feature Scaling

```

```

from sklearn.preprocessing import StandardScaler

```

```

sc = StandardScaler()

```

```

X_train = sc.fit_transform(X_train)

```

```

X_test = sc.transform(X_test)

```

```

results_4 = []; names = []; seed=42

```

```

for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_4 = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_4.append(cv_results_4)
    names.append(name)
    print(f"{name}, {cv_results_4.mean()}, {cv_results_4.std()}))")

```

```

LR, 0.7668103448275861, 0.08091636753181289))
LDA, 0.7652709359605911, 0.07755140356317015))
KNN, 0.7690270935960591, 0.0745836590047773))
CART, 0.7551724137931034, 0.07015987262294252))
NB, 0.7598522167487685, 0.07873013142757428))
SVM, 0.7686576354679804, 0.08000877623455863))
XGB, 0.7896551724137931, 0.07607369241829157))

```

5.5 5. Using forward Filling

```

[:]: df= pd.read_csv('/content/data/pima-indians-diabetes.csv', header= None)
df.columns= ['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age',
→'class']
df[df.columns[1:6]]= df[df.columns[1:6]].replace(0, np.nan)

df1 = df.loc[df['class'] == 1]
df2 = df.loc[df['class'] == 0]

```

```

[:]: df1.fillna(method= 'ffill', inplace=True)
df2.fillna(method= 'ffill', inplace=True)

df1.fillna(df1.mean(), inplace=True)
df2.fillna(df2.mean(), inplace=True)

dataframe = [df1, df2]
df = pd.concat(dataframe)

```

```

[:]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
all_inputs.shape

```

```

[:]: (768, 8)

```

```

[:]: (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    test_size=0.25,
    →random_state=1)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

```

```

results_5 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_5 = cross_val_score(model, X_train, y_train, cv=kfold,
    ↪scoring='accuracy')
    results_5.append(cv_results_5)
    names.append(name)
    print(f"{name}, {cv_results_5.mean()}, {cv_results_5.std()}")

```

```

LR, 0.7652093596059113, 0.07235394610657282))
LDA, 0.7669334975369458, 0.07328349403480619))
KNN, 0.7568965517241378, 0.0940046500259637))
CART, 0.7362068965517241, 0.053470972003503356))
NB, 0.7461206896551724, 0.06380840654648952))
SVM, 0.7600985221674877, 0.06766521362701008))
XGB, 0.7846059113300492, 0.05681452407049272))

```

5.6 Classification Comparison of Fill NaN Methods

```
[ ]: classification_comparison= pd.DataFrame(index=[i for i in names])
```

```
[ ]: names
```

```
[ ]: ['LR', 'LDA', 'KNN', 'CART', 'NB', 'SVM', 'XGB']
```

```

[ ]: dfresults = pd.DataFrame(results_1)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy']= df_mean
classification_comparison['sd']= df_sd

```

```

[ ]: dfresults = pd.DataFrame(results_2)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:

```

```

    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy_2']= df_mean
classification_comparison['sd_2']= df_sd

```

```

[:]: dfresults = pd.DataFrame(results_3)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy_3']= df_mean
classification_comparison['sd_3']= df_sd

```

```

[:]: dfresults = pd.DataFrame(results_4)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy_4']= df_mean
classification_comparison['sd_4']= df_sd

```

```

[:]: dfresults = pd.DataFrame(results_5)
dfresults=dfresults.T
dfresults.columns=names

```

```

df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy_5']= df_mean
classification_comparison['sd_5']= df_sd

```

```
[ ]: values= ['Acc', 'SD']
```

```
[ ]: Fill_NaN_Methods=['1_Remove', '2_Mean_fill', '3_Median_fill',
    '4_Back_fill', '5_Forwd_fill']
```

```
[ ]: idx = pd.MultiIndex.from_product([Fill_NaN_Methods, values],
    names=['Fill_NaN_Methods', 'values'])
```

```

classification_comparison.columns = idx
classification_comparison

```

```
[ ]: Fill_NaN_Methods 1_Remove      2_Mean_fill  ... 4_Back_fill 5_Forwd_fill
values              Acc      SD              Acc  ...              SD              Acc
SD
LR                  0.786  0.132      0.774  ...      0.083      0.765
0.074
LDA                  0.790  0.131      0.774  ...      0.080      0.767
0.075
KNN                  0.762  0.116      0.817  ...      0.077      0.757
0.096
CART                  0.690  0.134      0.880  ...      0.072      0.736
0.055
NB                   0.780  0.115      0.772  ...      0.081      0.746
0.065
SVM                  0.755  0.119      0.835  ...      0.082      0.760
0.069
XGB                  0.782  0.116      0.887  ...      0.078      0.785
0.058
```

```
[7 rows x 10 columns]
```

```
[ ]: classification_comparison.to_csv('Classification Comparison of Fill_NaN_Methods.
    ↳csv', index= False)
```

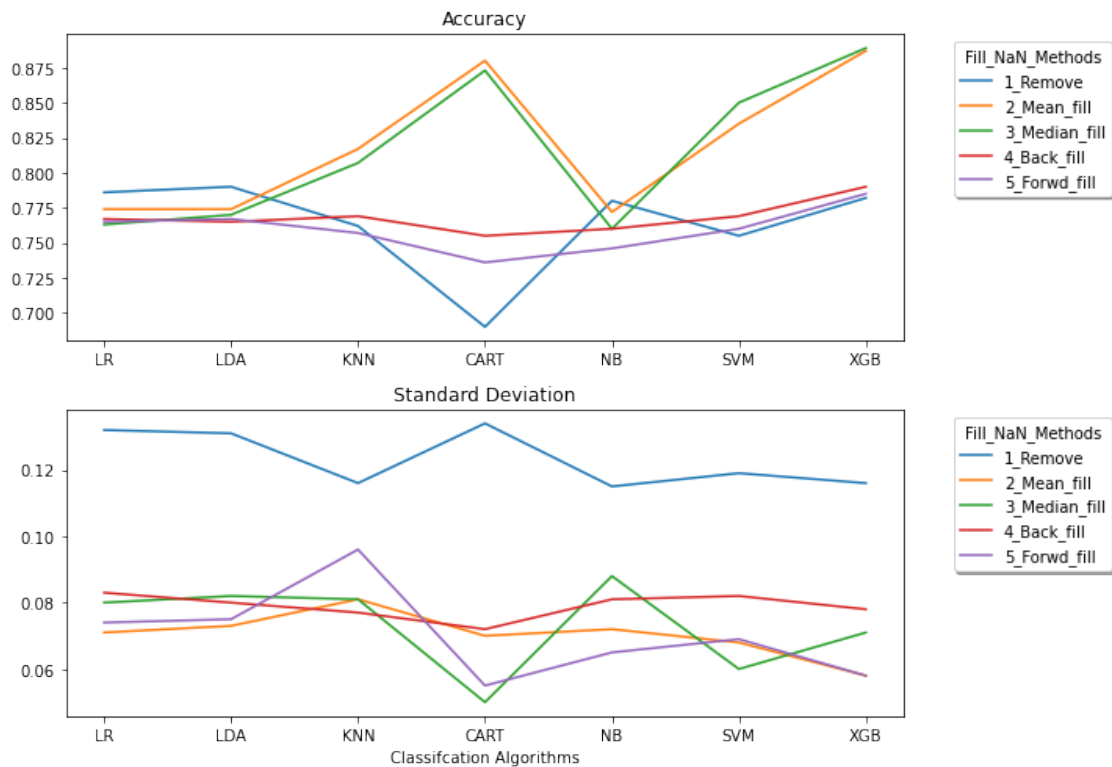
```
[ ]: plt.figure(figsize=(10,7))
    ax = plt.subplot(211)
```



```

classification_comparison.xs('Acc', axis=1, level='values').plot( ax=ax)
plt.title("Accuracy")
plt.legend(shadow=True, frameon=True, fancybox=True, title='Fill_NaN_Methods',
    ↳ bbox_to_anchor=(1.05, 1), loc='upper left')
ax = plt.subplot(212)
classification_comparison.xs('SD', axis=1, level='values').plot( ax=ax)
plt.title("Standard Deviation")
plt.xlabel("Classification Algorithms")
plt.legend(shadow=True, frameon=True, fancybox=True, title='Fill_NaN_Methods',
    ↳ bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()

```



```

[ ]: #classification_comparison= classification_comparison.drop('1_Remove', axis=1,
    ↳ level=0)

```

```

[ ]: classification_comparison.T.max()

```

```

[ ]: LR      0.786
     LDA      0.790
     KNN      0.817
     CART     0.880
     NB       0.780
     SVM      0.850
     XGB      0.889

```

dtype: float64

```
[ ]: Fill_NaN= classification_comparison.T.max()
```

```
[ ]: classification_comparison.T['CART']
```

```
[ ]: Fill_NaN_Methods  values
1_Remove             Acc      0.690
                     SD      0.134
2_Mean_fill          Acc      0.880
                     SD      0.070
3_Median_fill        Acc      0.873
                     SD      0.050
4_Back_fill          Acc      0.755
                     SD      0.072
5_Forwd_fill         Acc      0.736
                     SD      0.055
```

Name: CART, dtype: float64

```
[ ]: classification_comparison.T['XGB']
```

```
[ ]: Fill_NaN_Methods  values
1_Remove             Acc      0.782
                     SD      0.116
2_Mean_fill          Acc      0.887
                     SD      0.058
3_Median_fill        Acc      0.889
                     SD      0.071
4_Back_fill          Acc      0.790
                     SD      0.078
5_Forwd_fill         Acc      0.785
                     SD      0.058
```

Name: XGB, dtype: float64

```
[ ]: methods_comparison= pd.DataFrame(Fill_NaN, columns=['Fill_NaN'])
methods_comparison
```

```
[ ]:      Fill_NaN
LR      0.786
LDA      0.790
KNN      0.817
CART     0.880
NB      0.780
SVM      0.850
XGB      0.889
```

6 Selected Fill NaN Methods

The best fill NaN method is **Median_fill**

```
[ ]: df= pd.read_csv('/content/data/pima-indians-diabetes.csv', header= None)
df.columns= ['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age', 'class']
df[df.columns[1:6]]= df[df.columns[1:6]].replace(0, np.nan)

df1 = df.loc[df['class'] == 1]
df2 = df.loc[df['class'] == 0]

[ ]: df1.fillna(df1.median(), inplace=True)
df2.fillna(df2.median(), inplace=True)

dataframe = [df1, df2]
df = pd.concat(dataframe)
```

7 2. Add features

7.1 1. Add features based on BMI classification table

BMI classification table

- BMI CLASSIFICATION - > 30 :Obese - 25-30 :Overweight - 20-25 :Healthy weight range - 20-18 :Underweight - <18 :Very Underweight

```
[ ]: Obese= []
for i in df['BMI']:
    Obese.append(1 if i>30 else 0)
pd.value_counts(Obese)
```

```
[ ]: 1    476
0     292
dtype: int64
```

```
[ ]: df['Obese']= Obese
```

```
[ ]: Overweight= []
for i in df['BMI']:
    Overweight.append(1 if 30>i>25 else 0)
pd.value_counts(Overweight)
```

```
[ ]: 0    595
1     173
dtype: int64
```

```
[ ]: df['Overweight']= Overweight
```

```
[ ]: Healthy_weight= []
for i in df['BMI']:
    Healthy_weight.append(1 if 25>i>20 else 0)
pd.value_counts(Healthy_weight)
```

```
[ ]: 0    676
1     92
```

```

dtype: int64
[ ]: df['Healthy_weight']= Healthy_weight
[ ]: Underweight= []
    for i in df['BMI']:
        Underweight.append(1 if 20>i>18 else 0)
    pd.value_counts(Underweight)
[ ]: 0    755
     1     13
     dtype: int64
[ ]: df['Underweight']= Underweight
[ ]: Very_Underweight= []
    for i in df['BMI']:
        Very_Underweight.append(1 if i<18 else 0)
    pd.value_counts(Very_Underweight)
[ ]: 0    768
     dtype: int64
[ ]: #df['Very_Underweight']= Very_Underweight

```

7.2 2. Add features based on 2-Hour serum insulin

2-Hour serum insulin classification table

- insulin CLASSIFICATION - >140 :Normal - 140-199 :pre-diabetic - < 199 :diabetic

```

[ ]: Normal= []
    for i in df['insu']:
        Normal.append(1 if i<140 else 0)
    pd.value_counts(Normal)
[ ]: 1    457
     0    311
     dtype: int64
[ ]: df['Normal']= Normal
[ ]: pre_diabetic= []
    for i in df['insu']:
        pre_diabetic.append(1 if 199>i>140 else 0)
    pd.value_counts(pre_diabetic)
[ ]: 0    693
     1     75
     dtype: int64
[ ]: df['pre_diabetic']= pre_diabetic
[ ]: diabetic= []
    for i in df['insu']:

```

```
diabetic.append(1 if i>199 else 0)
pd.value_counts(diabetic)
```

```
[ ]: 0    541
      1    227
      dtype: int64
```

```
[ ]: df['diabetic']= diabetic
```

7.3 3. Add features based on Plasma glucose 2-Hour in an oral glucose tolerance test

2-Hour in an oral glucose classification table

- insulin CLASSIFICATION - >100 :Normal - 100-125 :pre_diabetic - < 125 :diabetic

```
[ ]: Normal_p= []
      for i in df['plas']:
          Normal_p.append(1 if i<100 else 0)
      pd.value_counts(Normal_p)
```

```
[ ]: 0    576
      1    192
      dtype: int64
```

```
[ ]: df['Normal_p']= Normal_p
```

```
[ ]: pre_diabetic_p= []
      for i in df['plas']:
          pre_diabetic_p.append(1 if 125>i>100 else 0)
      pd.value_counts(pre_diabetic_p)
```

```
[ ]: 0    522
      1    246
      dtype: int64
```

```
[ ]: df['pre_diabetic_p']= pre_diabetic_p
```

```
[ ]: diabetic_p= []
      for i in df['plas']:
          diabetic_p.append(1 if i>125 else 0)
      pd.value_counts(diabetic_p)
```

```
[ ]: 0    469
      1    299
      dtype: int64
```

```
[ ]: df['diabetic_p']= diabetic_p
```

```
[ ]: df= df[['preg', 'plas', 'pres_mm', 'skin_mm', 'insu', 'BMI', 'DPF', 'age',
             'Obese', 'Overweight', 'Healthy_weight', 'Underweight',
             'Normal', 'pre_diabetic', 'diabetic', 'Normal_p',
             'pre_diabetic_p', 'diabetic_p', 'class']]
```

```
[ ]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
all_inputs.shape
```

```
[ ]: (768, 18)
```

7.4 Classification Comparison

```
[ ]: results_1 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_1 = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_1.append(cv_results_1)
    names.append(name)
    print(f"{name}, {cv_results_1.mean()}, {cv_results_1.std()}")
```

```
LR, 0.7652093596059113, 0.07235394610657282))
LDA, 0.7669334975369458, 0.07328349403480619))
KNN, 0.7568965517241378, 0.0940046500259637))
CART, 0.724076354679803, 0.06990961268480901))
NB, 0.7461206896551724, 0.06380840654648952))
SVM, 0.7600985221674877, 0.06766521362701008))
XGB, 0.7846059113300492, 0.05681452407049272))
```

```
[ ]: classification_comparison= pd.DataFrame(index=[i for i in names])
```

```
[ ]: dfresults = pd.DataFrame(results_1)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy']= df_mean
classification_comparison['sd']= df_sd
classification_comparison
```

```
[ ]:      Accuracy      sd
LR      0.765    0.074
LDA      0.767    0.075
```

KNN	0.757	0.096
CART	0.724	0.072
NB	0.746	0.065
SVM	0.760	0.069
XGB	0.785	0.058

```
[ ]: Add_F= classification_comparison.T.max()
```

```
[ ]: methods_comparison['Add_F']= Add_F
methods_comparison
```

```
[ ]:      Fill_NaN  Add_F
LR      0.786  0.765
LDA      0.790  0.767
KNN      0.817  0.757
CART      0.880  0.724
NB      0.780  0.746
SVM      0.850  0.760
XGB      0.889  0.785
```

8 3. Automatic Outlier Detection Algorithms Comparison

8.1 1. DBSCAN

```
[ ]: from sklearn.cluster import DBSCAN
def remove_outliers_DBSCAN(df,eps,min_samples):
    outlier_detection = DBSCAN(eps = eps, min_samples = min_samples)
    clusters = outlier_detection.fit_predict(df.values.reshape(-1,1))
    data = pd.DataFrame()
    data['cluster'] = clusters
    return data['cluster']
```

```
[ ]: outlier_index1=[]
for col in df.columns[0:-1]:
    clusters=remove_outliers_DBSCAN((df[col]), .2, 2)
    df_cluster=pd.DataFrame(clusters)
    outlier_index1= outlier_index1+(list(df_cluster.
    →index[df_cluster['cluster']==-1]))
outlier_index1=list(set(outlier_index1))
print(len(outlier_index1))

DBSCAN_df=df.drop(outlier_index1)
#DBSCAN_df.to_csv('DBSCAN_df.csv', index= False)

#!mkdir Outlier_Detection_DFs
#!mv DBSCAN_df.csv Outlier_Detection_DFs/

all_inputs = DBSCAN_df[DBSCAN_df.columns[0:-1]].values
```

```

all_labels = DBSCAN_df['class'].values

(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_1 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_1 = cross_val_score(model, X_train, y_train, cv=kfold,
    ↳scoring='accuracy')
    results_1.append(cv_results_1)
    names.append(name)
    print(f"{name}, {cv_results_1.mean()}, {cv_results_1.std()}")

```

```

136
LR, 0.8817934782608697, 0.07055194255956952))
LDA, 0.879981884057971, 0.06803706068962845))
KNN, 0.8903985507246377, 0.06055035130416818))
CART, 0.8694746376811594, 0.08009925407624391))
NB, 0.6959239130434783, 0.09377738870386135))
SVM, 0.8842391304347826, 0.062361745847078154))
XGB, 0.8971014492753623, 0.06704269346708851))

```

8.2 2. Isolation Forest

```

[ ]: to_model_columns=df.columns[0:-1]
from sklearn.ensemble import IsolationForest
clf=IsolationForest(n_estimators=99, max_samples='auto',
    contamination=0.2,
    max_features=1.0 , bootstrap=False, n_jobs=-1,
    random_state=42, verbose=0 )

clf.fit(df[to_model_columns])
pred = clf.predict(df[to_model_columns])
df['anomaly']=pred
outliers=df.loc[df['anomaly']==-1]
outlier_index2=list(outliers.index)
print(outlier_index2)
#Find the number of anomalies and normal points here points classified -1 are
↳anomalous
print(df['anomaly'].value_counts())
df.drop('anomaly', axis='columns', inplace=True)

```



```

Isolation_Forest_df=df.drop(outlier_index2)
Isolation_Forest_df['class'].value_counts()
#Isolation_Forest_df.to_csv('Isolation_Forest_df.csv', index= False)
#!mv Isolation_Forest_df.csv Outlier_Detection_DFs/
all_inputs = Isolation_Forest_df[Isolation_Forest_df.columns[0:-1]].values
all_labels = Isolation_Forest_df['class'].values
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_2 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_2 = cross_val_score(model, X_train, y_train, cv=kfold,
    ↳scoring='accuracy')
    results_2.append(cv_results_2)
    names.append(name)
    print(f"{name}, {cv_results_2.mean()}, {cv_results_2.std()}")

```

```

[2, 4, 8, 13, 14, 24, 38, 39, 43, 93, 99, 114, 120, 125, 128, 129, 130, 131,
152, 159, 165, 177, 214, 218, 236, 243, 245, 254, 259, 270, 276, 284, 292, 293,
296, 298, 308, 319, 322, 323, 356, 370, 400, 408, 429, 444, 445, 476, 498, 510,
516, 539, 542, 579, 584, 588, 590, 595, 614, 618, 646, 647, 676, 689, 709, 719,
731, 740, 749, 12, 27, 28, 33, 35, 50, 51, 57, 59, 68, 90, 91, 94, 105, 106,
194, 204, 223, 228, 239, 247, 248, 258, 279, 294, 307, 311, 316, 320, 325, 330,
333, 374, 379, 382, 385, 390, 392, 395, 396, 405, 418, 420, 428, 434, 438, 452,
453, 459, 460, 464, 466, 489, 499, 507, 511, 512, 519, 526, 528, 537, 549, 582,
593, 597, 607, 609, 610, 617, 621, 633, 639, 644, 657, 668, 672, 673, 679, 688,
713, 717, 718, 733, 738, 763]
1      614
-1     154
Name: anomaly, dtype: int64
LR, 0.8739130434782607, 0.06860753842634564))
LDA, 0.8760869565217391, 0.06341718324267313))
KNN, 0.8782608695652174, 0.0507039295203939))
CART, 0.8804347826086956, 0.04534055133611785))
NB, 0.8413043478260869, 0.07810975749226816))
SVM, 0.8739130434782607, 0.059772726455945724))
XGB, 0.8956521739130435, 0.055679341195068247))

```

8.3 3. Minimum Covariance Determinant

```
[ ]: to_model_columns=df.columns[:-1]
from sklearn.covariance import EllipticEnvelope
ee = EllipticEnvelope(contamination=.01, )
ee.fit(df[to_model_columns])
pred = ee.predict(df[to_model_columns])
df['anomaly']=pred
outliers=df.loc[df['anomaly']==-1]
outlier_index3=list(outliers.index)
print(outlier_index3)
#Find the number of anomalies and normal points here points classified -1 are
    →anomalous
print(df['anomaly'].value_counts())
df.drop('anomaly', axis='columns', inplace=True)
Minimum_Covariance_Determinant_df=df.drop(outlier_index3)
Minimum_Covariance_Determinant_df['class'].value_counts()
#Minimum_Covariance_Determinant_df.to_csv('Minimum_Covariance_Determinant_df.
    →csv', index= False)
#!mv Minimum_Covariance_Determinant_df.csv Outlier_Detection_DFs/
all_inputs =
    →Minimum_Covariance_Determinant_df[Minimum_Covariance_Determinant_df.
    →columns[0:-1]].values
all_labels = Minimum_Covariance_Determinant_df['class'].values
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    →test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_3 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_3 = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_3.append(cv_results_3)
    names.append(name)
    print(f"{name}, {cv_results_3.mean()}, {cv_results_3.std()}))")
```

```
[15, 245, 62, 217, 453, 519, 639, 745]
1      760
-1      8
Name: anomaly, dtype: int64
LR, 0.8751847290640395, 0.061953616476605494))
LDA, 0.875, 0.06923165135030956))
```

```
KNN, 0.862807881773399, 0.06063992388560086))
CART, 0.8504310344827586, 0.06325360413847252))
NB, 0.6526477832512315, 0.08628908379926371))
SVM, 0.8573275862068966, 0.07557531190421447))
XGB, 0.8858990147783251, 0.04310569135613251))
```

8.4 4. Local Outlier Factor

```
[ ]: to_model_columns=df.columns[0:-1]
from sklearn.neighbors import LocalOutlierFactor
lof = LocalOutlierFactor(novelty=True, n_jobs=1,
                        n_neighbors=5, contamination=0.15, leaf_size= 60 )
lof.fit(df[to_model_columns])
pred = lof.predict(df[to_model_columns])
df['anomaly']=pred
outliers=df.loc[df['anomaly']==-1]
outlier_index4=list(outliers.index)
print(outlier_index4)
#Find the number of anomalies and normal points here points classified -1 are
→anomalous
print(df['anomaly'].value_counts())
df.drop('anomaly', axis='columns', inplace=True)
Local_Outlier_Factor_df=df.drop(outlier_index4)
Local_Outlier_Factor_df['class'].value_counts()
#Local_Outlier_Factor_df.to_csv('Local_Outlier_Factor_df.csv', index= False)
#!mv Local_Outlier_Factor_df.csv Outlier_Detection_DFs/
all_inputs = Local_Outlier_Factor_df[Local_Outlier_Factor_df.columns[0:-1]].
    →values
all_labels = Local_Outlier_Factor_df['class'].values
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    →test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_4 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_4 = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_4.append(cv_results_4)
    names.append(name)
    print(f"{name}, {cv_results_4.mean()}, {cv_results_4.std()}")
```

```
[13, 39, 43, 120, 125, 132, 177, 187, 193, 227, 237, 238, 242, 293, 328, 440,
444, 445, 476, 545, 579, 595, 647, 691, 693, 716, 740, 759, 18, 28, 33, 55, 57,
62, 86, 95, 106, 117, 135, 147, 150, 182, 223, 228, 244, 247, 250, 256, 277,
307, 320, 346, 362, 379, 396, 434, 456, 459, 460, 466, 511, 519, 537, 573, 575,
597, 653, 657, 658, 672, 680, 763]
1      696
-1      72
Name: anomaly, dtype: int64
LR, 0.8618945868945869, 0.07896991242112261))
LDA, 0.8618233618233617, 0.07821148501142053))
KNN, 0.8580484330484331, 0.0758341665289647))
CART, 0.8619658119658119, 0.06257745746129312))
NB, 0.6515669515669515, 0.11241599903766386))
SVM, 0.8752849002849002, 0.056754316011546))
XGB, 0.904059829059829, 0.07256471637788393))
```

8.5 5. One-Class SVM

```
[ ]: to_model_columns=df.columns[0:-1]
from sklearn.svm import OneClassSVM

ocs = OneClassSVM(nu=0.15, )
ocs.fit(df[to_model_columns])
pred = ocs.predict(df[to_model_columns])
df['anomaly']=pred
outliers=df.loc[df['anomaly']==-1]
outlier_index5=list(outliers.index)
print(outlier_index5)

#Find the number of anomalies and normal points here points classified -1 are
→anomalous
print(df['anomaly'].value_counts())
df.drop('anomaly', axis='columns', inplace=True)
One_Class_SVM_df=df.drop(outlier_index5)
One_Class_SVM_df['class'].value_counts()

#One_Class_SVM_df.to_csv('One_Class_SVM_df.csv', index= False)
#!mv One_Class_SVM_df.csv Outlier_Detection_DFs/

all_inputs = One_Class_SVM_df[One_Class_SVM_df.columns[0:-1]].values
all_labels = One_Class_SVM_df['class'].values
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
→test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```

X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_5 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_5 = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_5.append(cv_results_5)
    names.append(name)
    print(f"{name}, {cv_results_5.mean()}, {cv_results_5.std()}")

```

```

[8, 13, 22, 43, 53, 56, 109, 111, 186, 199, 206, 220, 231, 254, 296, 319, 323,
359, 360, 370, 375, 388, 408, 409, 415, 425, 440, 445, 480, 498, 506, 561, 579,
584, 606, 612, 655, 661, 695, 715, 753, 759, 32, 40, 51, 52, 54, 62, 68, 73, 92,
103, 108, 112, 139, 144, 153, 162, 173, 174, 182, 203, 225, 228, 232, 234, 247,
248, 258, 260, 273, 279, 286, 288, 290, 302, 353, 364, 392, 395, 412, 441, 459,
462, 466, 482, 486, 487, 489, 519, 520, 534, 537, 548, 549, 553, 566, 572, 574,
597, 607, 617, 625, 639, 645, 672, 679, 680, 707, 710, 711, 713, 747, 760]
1      654
-1     114
Name: anomaly, dtype: int64
LR, 0.8913333333333334, 0.06084954121985363))
LDA, 0.8955, 0.05515407711332157))
KNN, 0.8874166666666667, 0.06921077910589103))
CART, 0.8877499999999999, 0.0630878466020904))
NB, 0.6793333333333333, 0.08945110396188523))
SVM, 0.8873333333333333, 0.06971250485625469))
XGB, 0.8915, 0.058790825058941905))

```

8.6 Comparison

```

[: classification_comparison= pd.DataFrame(index=[i for i in names])
[: dfresults = pd.DataFrame(results_1)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

```

```
classification_comparison['Accuracy_1']= df_mean
classification_comparison['sd_1']= df_sd
```

```
[ ]: dfresults = pd.DataFrame(results_2)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy_2']= df_mean
classification_comparison['sd_2']= df_sd
```

```
[ ]: dfresults = pd.DataFrame(results_3)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy_3']= df_mean
classification_comparison['sd_3']= df_sd
```

```
[ ]: dfresults = pd.DataFrame(results_4)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
```

```

    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy_4']= df_mean
classification_comparison['sd_4']= df_sd

```

```

[:]: dfresults = pd.DataFrame(results_5)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 3) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 3) for num in df_sd]

classification_comparison['Accuracy_5']= df_mean
classification_comparison['sd_5']= df_sd

```

```

[:]: classification_comparison

```

```

[:]:
      Accuracy_1  sd_1  Accuracy_2  ...  sd_4  Accuracy_5  sd_5
LR              0.882  0.072      0.874  ...  0.081      0.891  0.062
LDA              0.880  0.070      0.876  ...  0.080      0.895  0.057
KNN              0.890  0.062      0.878  ...  0.078      0.887  0.071
CART              0.869  0.082      0.880  ...  0.064      0.888  0.065
NB               0.696  0.096      0.841  ...  0.115      0.679  0.092
SVM              0.884  0.064      0.874  ...  0.058      0.887  0.072
XGB              0.897  0.069      0.896  ...  0.074      0.892  0.060

```

[7 rows x 10 columns]

```

[:]: values= ['Acc', 'SD']

```

```

[:]: DFs=[ '1_DBSCAN_df', '2_IF_df', '3_MCD_df', '4_LOF_df', '5_OCSVM_df']

```

```

[:]: idx = pd.MultiIndex.from_product([DFs, values],
                                     names=['DFs', 'values'])

classification_comparison.columns = idx
classification_comparison

```

```

[:]: DFs      1_DBSCAN_df      2_IF_df      ...  4_LOF_df      5_OCSVM_df
values      Acc      SD      Acc      SD      ...      Acc      SD      Acc      SD
LR          0.882  0.072  0.874  0.070  ...  0.862  0.081  0.891  0.062


```

LDA	0.880	0.070	0.876	0.065	...	0.862	0.080	0.895	0.057
KNN	0.890	0.062	0.878	0.052	...	0.858	0.078	0.887	0.071
CART	0.869	0.082	0.880	0.047	...	0.862	0.064	0.888	0.065
NB	0.696	0.096	0.841	0.080	...	0.652	0.115	0.679	0.092
SVM	0.884	0.064	0.874	0.061	...	0.875	0.058	0.887	0.072
XGB	0.897	0.069	0.896	0.057	...	0.904	0.074	0.892	0.060

[7 rows x 10 columns]

```
[ ]: #classification_comparison.to_csv('classification_comparison.csv', index=
    ↳False)
```

```
[ ]: plt.figure(figsize=(10,7))
ax = plt.subplot(211)
classification_comparison.xs('Acc', axis=1, level='values').plot( ax=ax)
plt.title("Accuracy")
plt.legend(shadow=True, frameon=True, fancybox=True, title='DFs',
    ↳bbox_to_anchor=(1.05, 1), loc='upper left')
ax = plt.subplot(212)
classification_comparison.xs('SD', axis=1, level='values').plot( ax=ax)
plt.title("Standard Deviation")
plt.xlabel("Classification Algorithms")
plt.legend(shadow=True, frameon=True, fancybox=True, title='DFs',
    ↳bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
```




```
[ ]: classification_comparison.T.max()
```

```
[ ]: LR      0.891  
LDA      0.895  
KNN      0.890  
CART      0.888  
NB       0.841  
SVM      0.887  
XGB      0.904  
dtype: float64
```

```
[ ]: classification_comparison.T['CART']
```

```
[ ]: DFs      values  
1_DBSCAN_df  Acc      0.869  
              SD      0.082  
2_IF_df      Acc      0.880  
              SD      0.047  
3_MCD_df     Acc      0.850  
              SD      0.065  
4_LOF_df     Acc      0.862  
              SD      0.064  
5_OCSVM_df   Acc      0.888  
              SD      0.065  
Name: CART, dtype: float64
```

```
[ ]: classification_comparison.T['XGB']
```

```
[ ]: DFs      values  
1_DBSCAN_df  Acc      0.897  
              SD      0.069  
2_IF_df      Acc      0.896  
              SD      0.057  
3_MCD_df     Acc      0.886  
              SD      0.044  
4_LOF_df     Acc      0.904  
              SD      0.074  
5_OCSVM_df   Acc      0.892  
              SD      0.060  
Name: XGB, dtype: float64
```

```
[ ]: Outlier=classification_comparison.T.max()
```

```
[ ]: methods_comparison['Outlier']= Outlier  
methods_comparison
```

```
[ ]:      Fill_NaN  Add_F  Outlier  
LR      0.786  0.765   0.891  
LDA      0.790  0.767   0.895
```

KNN	0.817	0.757	0.890
CART	0.880	0.724	0.888
NB	0.780	0.746	0.841
SVM	0.850	0.760	0.887
XGB	0.889	0.785	0.904

8.7 Selected Outlier Algorithm DF

The best Outlier Algorithm is Local Outlier Factor, so it performed on the data.

```
[ ]: df= Local_Outlier_Factor_df
```

```
[ ]: df['class'].value_counts()
```

```
[ ]: 0    456
      1    240
      Name: class, dtype: int64
```

9 4. Feature Selection Methods Comparison

9.1 1. Removing features with low variance

VarianceThreshold is a simple baseline approach to feature selection. It removes all features whose variance doesn't meet some threshold. By default, it removes all zero-variance features, i.e. features that have the same value in all samples.*italicised text*

```
[ ]: df.var().nlargest(20)
```

```
[ ]: insu            6739.204056
      plas           882.869127
      age           129.165945
      pres_mm       118.224352
      skin_mm        64.182265
      BMI            42.363108
      preg           11.016919
      Normal         0.239938
      Obese          0.236796
      diabetic_p     0.235070
      class          0.226247
      pre_diabetic_p 0.220590
      diabetic       0.207492
      Normal_p       0.189909
      Overweight     0.178068
      Healthy_weight 0.104085
      DPF            0.100794
      pre_diabetic   0.084791
      Underweight    0.015577
      dtype: float64
```

```

[:]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
print(all_inputs.shape)

from sklearn.feature_selection import VarianceThreshold
sel = VarianceThreshold(threshold= 0.104085)
all_inputs = sel.fit_transform(all_inputs)
print(all_inputs.shape)
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
→test_size=0.25, random_state=1)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_1 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_1 = cross_val_score(model, X_train, y_train, cv=kfold,
→scoring='accuracy')
    results_1.append(cv_results_1)
    names.append(name)
    print(f"{name}, {cv_results_1.mean()}, {cv_results_1.std()}")

```

```

(696, 18)
(696, 14)
LR, 0.862037037037037, 0.06597345014640185))
LDA, 0.8618945868945869, 0.07555832044504261))
KNN, 0.8582621082621082, 0.0759855017679339))
CART, 0.8413105413105415, 0.07110445627784064))
NB, 0.8447293447293449, 0.06945960094664742))
SVM, 0.8773504273504272, 0.08218138723352486))
XGB, 0.8984330484330485, 0.06469354046681688))

```

Univariate Feature Selection

9.2 2. UFS SelectKBest Select K Best

it removes all but the highest scoring features

```

[:]: all_inputs = df[df.columns[0:-1]].values
from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler()
all_inputs = sc.fit_transform(all_inputs)
all_labels = df['class'].values
print(all_inputs.shape)

```

```

from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2

all_inputs = SelectKBest(chi2, k=15).fit_transform(all_inputs, all_labels)
print(all_inputs.shape)
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
→test_size=0.25, random_state=1)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_2 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_2 = cross_val_score(model, X_train, y_train, cv=kfold,
→scoring='accuracy')
    results_2.append(cv_results_2)
    names.append(name)
    print(f"{name}, {cv_results_2.mean()}, {cv_results_2.std()}")

```

```

(696, 18)
(696, 15)
LR, 0.8619658119658119, 0.06138411185508111))
LDA, 0.8734330484330485, 0.06463105191758405))
KNN, 0.8560541310541311, 0.07311658592755686))
CART, 0.861894586894587, 0.06727278674035306))
NB, 0.6435185185185184, 0.08472835202709286))
SVM, 0.8657407407407408, 0.07563804110411428))
XGB, 0.9041310541310541, 0.06387280762652076))

```

9.3 3. UFSSelectFpr' False Positive Rate test.

Filter: Select the p values below alpha based on a FPR test. a smaller p-value bears more significance as it can tell you that the hypothesis may not explain the observation fairly. If one or more of these probabilities turn out to be less than or equal to , the level of significance, we reject the null hypothesis. For a true null hypothesis, p can take on any value between 0 and 1 with equal likeliness. For a true alternative hypothesis, p-values likely fall closer to 0.

```

[ ]: all_inputs = df[df.columns[0:-1]].values
sc = MinMaxScaler()
all_inputs = sc.fit_transform(all_inputs)
all_labels = df['class'].values
from sklearn.feature_selection import SelectFpr, chi2
all_inputs = SelectFpr(chi2, alpha= 0.09
→).fit_transform(all_inputs, all_labels)

```

```

print(all_inputs.shape)
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳test_size=0.25, random_state=1)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_3 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_3 = cross_val_score(model, X_train, y_train, cv=kfold,
    ↳scoring='accuracy')
    results_3.append(cv_results_3)
    names.append(name)
    print(f"{name}, {cv_results_3.mean()}, {cv_results_3.std()}")

```

```

(696, 16)
LR, 0.8677350427350428, 0.05641326900643618))
LDA, 0.8733618233618234, 0.06583277504291998))
KNN, 0.8447293447293449, 0.06729619733314672))
CART, 0.8527065527065527, 0.07052418683031295))
NB, 0.6684472934472934, 0.08571531178496568))
SVM, 0.8753561253561255, 0.07497872415709027))
XGB, 0.8983618233618234, 0.06592114099112054))

```

9.4 4. Feature selection using SelectFromModel

9.4.1 L1-based feature selection

Linear models penalized with the L1 norm have sparse solutions: many of their estimated coefficients are zero. When the goal is to reduce the dimensionality of the data to use with another classifier, they can be used along with `SelectFromModel` to select the non-zero coefficients. In particular, sparse estimators useful for this purpose are the Lasso for regression, and of LogisticRegression and LinearSVC for classification:

```

[ ]: all_inputs = df[df.columns[0:-1]].values
sc = StandardScaler()
all_inputs = sc.fit_transform(all_inputs)
all_labels = df['class'].values

from sklearn.svm import LinearSVC
from sklearn.feature_selection import SelectFromModel

lsvc = LinearSVC(C=.09 , penalty="l1", dual=False).fit(all_inputs, all_labels)
model = SelectFromModel(lsvc, prefit=True)

```

```

all_inputs = model.transform(all_inputs)
print(all_inputs.shape)

from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳test_size=0.25, random_state=1)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_4 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_4 = cross_val_score(model, X_train, y_train, cv=kfold,
    ↳scoring='accuracy')
    results_4.append(cv_results_4)
    names.append(name)
    print(f"{name}, {cv_results_4.mean()}, {cv_results_4.std()}")

```

```

(696, 14)
LR, 0.8716524216524217, 0.057349054491092656))
LDA, 0.8753561253561255, 0.05958861601904569))
KNN, 0.8657407407407408, 0.06845151790492156))
CART, 0.8467948717948717, 0.06182643886679986))
NB, 0.8504985754985755, 0.07195561951720005))
SVM, 0.865883190883191, 0.07546399654292588))
XGB, 0.9061253561253562, 0.05633700530500054))

```

9.5 5. Sequential Feature Selection

(Selecting features based on importance) The features with the highest absolute coef_ value are considered the most important.

```

[:]: all_inputs = df[df.columns[0:-1]].values
sc = StandardScaler()
all_inputs = sc.fit_transform(all_inputs)
all_labels = df['class'].values
all_inputs.shape

```

```

[:]: (696, 18)

```

Feature importance from coefficients

```


[:]: from sklearn.linear_model import LassoCV
from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler()
all_inputs = sc.fit_transform(all_inputs)

```

```

lasso = LassoCV().fit(all_inputs, all_labels)
importance = np.abs(lasso.coef_)
feature_names = np.array(all_labels)
plt.bar(height=importance, x=df.columns[0:-1])
plt.xticks(rotation=80)
plt.title("Feature importances via coefficients")
plt.show()

```



```

[: np.sort(importance)

```

```

[: array([0.00516833, 0.01704968, 0.01878721, 0.03886262, 0.04583879,
0.05275003, 0.0706723 , 0.10114942, 0.10380875, 0.12037679,
0.1681799 , 0.19411578, 0.21222111, 0.23328783, 0.40594585,
0.51628021, 0.78462425, 0.85159204])

```

```

[: all_inputs = df[df.columns[0:-1]].values
sc = MinMaxScaler()
all_inputs = sc.fit_transform(all_inputs)
all_labels = df['class'].values
print(all_inputs.shape)

```

```

from sklearn.feature_selection import SelectFromModel

threshold = 0.03886262

sfm = SelectFromModel(lasso, threshold=threshold).fit(all_inputs, all_labels)
selected_Features = df.columns[:-1][sfm.get_support()]
print("Features selected by SelectFromModel: ",
      f"{df.columns[0:-1][sfm.get_support()]}")

all_inputs = df[selected_Features].values
all_labels = df['class'].values
print(all_inputs.shape)
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
→test_size=0.25, random_state=1)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_5 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_5 = cross_val_score(model, X_train, y_train, cv=kfold,
→scoring='accuracy')
    results_5.append(cv_results_5)
    names.append(name)
    print(f"{name}, {cv_results_5.mean()}, {cv_results_5.std()}")

```

(696, 18)

Features selected by SelectFromModel: Index(['preg', 'plas', 'skin_mm', 'insu',
'BMI', 'DPF', 'age', 'Obese',
'Healthy_weight', 'pre_diabetic', 'diabetic', 'Normal_p',
'pre_diabetic_p', 'diabetic_p'],
dtype='object')

(696, 14)

LR, 0.8735042735042736, 0.05453552591347405))
LDA, 0.8752849002849004, 0.06089218750812922))
KNN, 0.8428062678062679, 0.06207995349260721))
CART, 0.8544159544159544, 0.061490959824290425))
NB, 0.8141737891737891, 0.08181584589396765))
SVM, 0.8772079772079773, 0.07476067936357167))
XGB, 0.9061253561253559, 0.05890429657420025))

9.6 6. Principal Component Analysis

```
[ ]: all_inputs = df[df.columns[0:-1]].values
all_labels = df['class'].values
all_inputs.shape

[ ]: (696, 18)

[ ]: from sklearn.decomposition import PCA
pca = PCA(n_components= 12)
pca.fit(all_inputs)

[ ]: PCA(copy=True, iterated_power='auto', n_components=12, random_state=None,
      svd_solver='auto', tol=0.0, whiten=False)

[ ]: print(pca.explained_variance_)

[6.99890877e+03 6.54746066e+02 1.52460758e+02 9.32832577e+01
 6.02073986e+01 2.07172485e+01 7.41590680e+00 3.10281425e-01
 2.16565656e-01 1.61184212e-01 1.22365951e-01 9.54000208e-02]

[ ]: all_inputs = pca.transform(all_inputs)
all_inputs.shape
from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    →test_size=0.25, random_state=1)
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_6 = []
names = []
seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_6 = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_6.append(cv_results_6)
    names.append(name)
    print(f"{name}, {cv_results_6.mean()}, {cv_results_6.std()}"))

LR, 0.8620370370370372, 0.054964854020312))
LDA, 0.8619658119658121, 0.06498115844858557))
KNN, 0.8619658119658121, 0.0725124397045091))
CART, 0.817877492877493, 0.09005764588838602))
NB, 0.8486467236467238, 0.0670679642257839))
SVM, 0.8715099715099717, 0.07539045098468716))
XGB, 0.8563390313390314, 0.08044621919946748))
```

9.7 Classification Comparison of Feature Selection Methods

```
[ ]: classification_comparison= pd.DataFrame(index=[i for i in names])
```

```
[ ]: dfresults = pd.DataFrame(results_1)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy']= df_mean
classification_comparison['sd']= df_sd
```

```
[ ]: dfresults = pd.DataFrame(results_2)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_2']= df_mean
classification_comparison['sd_2']= df_sd
```

```
[ ]: dfresults = pd.DataFrame(results_3)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
```

```

    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_3']= df_mean
classification_comparison['sd_3']= df_sd

```

```

[: dfresults = pd.DataFrame(results_4)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_4']= df_mean
classification_comparison['sd_4']= df_sd

```

```

[: dfresults = pd.DataFrame(results_5)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_5']= df_mean
classification_comparison['sd_5']= df_sd

```

```

[: dfresults = pd.DataFrame(results_6)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)

```

```

df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_6']= df_mean
classification_comparison['sd_6']= df_sd

```

```
[ ]: values= ['Acc', 'SD']
```

```
[ ]: FS_Methods=['1_RFLVar', '2_UFSKBest', '3_UFSFpr',
                '4_FSU', '5_SFS', '6_PCA']
```

```
[ ]: idx = pd.MultiIndex.from_product([FS_Methods, values],
                                     names=['FS_Methods', 'values'])

classification_comparison.columns = idx
classification_comparison

```

```
[ ]: FS_Methods 1_RFLVar      2_UFSKBest      ...  5_SFS      6_PCA
values          Acc      SD          Acc      SD      ...      Acc      SD      Acc
SD
LR              0.8620  0.0677      0.8620  0.0630  ...  0.8735  0.0560  0.8620
0.0564
LDA              0.8619  0.0775      0.8734  0.0663  ...  0.8753  0.0625  0.8620
0.0667
KNN              0.8583  0.0780      0.8561  0.0750  ...  0.8428  0.0637  0.8620
0.0744
CART              0.8413  0.0730      0.8619  0.0690  ...  0.8544  0.0631  0.8179
0.0924
NB              0.8447  0.0713      0.6435  0.0869  ...  0.8142  0.0839  0.8486
0.0688
SVM              0.8774  0.0843      0.8657  0.0776  ...  0.8772  0.0767  0.8715
0.0773
XGB              0.8984  0.0664      0.9041  0.0655  ...  0.9061  0.0604  0.8563
0.0825
```

[7 rows x 12 columns]

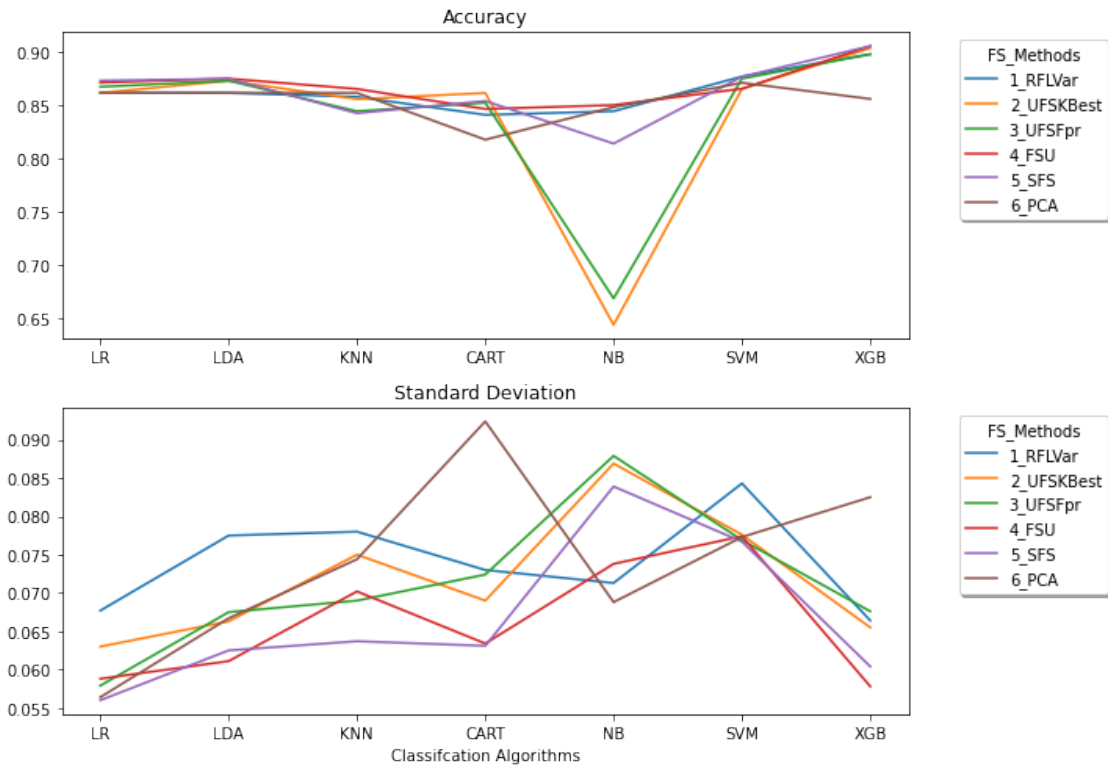
```
[ ]: classification_comparison.to_csv('Classification Comparison of Feature_
    ↳Selection Methods.csv', index= False)
```

```
[ ]: plt.figure(figsize=(10,7))
ax = plt.subplot(211)
classification_comparison.xs('Acc', axis=1, level='values').plot( ax=ax)
plt.title("Accuracy")
plt.legend(shadow=True, frameon=True, fancybox=True, title='FS_Methods',
    ↳bbox_to_anchor=(1.05, 1), loc='upper left')
ax = plt.subplot(212)
```

```

classification_comparison.xs('SD', axis=1, level='values').plot( ax=ax)
plt.title("Standard Deviation")
plt.xlabel("Classification Algorithms")
plt.legend(shadow=True, frameon=True, fancybox=True, title='FS_Methods',
           bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()

```



```

[: classification_comparison.T.max()

```

```

[: LR      0.8735
   LDA     0.8754
   KNN     0.8657
   CART    0.8619
   NB      0.8505
   SVM     0.8774
   XGB     0.9061
   dtype: float64

```

```

[: Feature =classification_comparison.T.max()

```

```

[: classification_comparison.T['XGB']

```

```

[: FS_Methods  values
   1_RFLVar    Acc      0.8984
               SD      0.0664

```

2_UFSKBest	Acc	0.9041
	SD	0.0655
3_UFSFpr	Acc	0.8984
	SD	0.0676
4_FSU	Acc	0.9061
	SD	0.0578
5_SFS	Acc	0.9061
	SD	0.0604
6_PCA	Acc	0.8563
	SD	0.0825

Name: XGB, dtype: float64

```
[ ]: classification_comparison.T['CART']
```

```
[ ]: FS_Methods values
1_RFLVar    Acc    0.8413
           SD    0.0730
2_UFSKBest  Acc    0.8619
           SD    0.0690
3_UFSFpr    Acc    0.8527
           SD    0.0724
4_FSU       Acc    0.8468
           SD    0.0634
5_SFS       Acc    0.8544
           SD    0.0631
6_PCA       Acc    0.8179
           SD    0.0924
```

Name: CART, dtype: float64

```
[ ]: methods_comparison['Feature']= Feature
methods_comparison
```

```
[ ]:      Fill_NaN  Add_F  Outlier  Feature
LR      0.786  0.765    0.891   0.8735
LDA      0.790  0.767    0.895   0.8754
KNN      0.817  0.757    0.890   0.8657
CART     0.880  0.724    0.888   0.8619
NB       0.780  0.746    0.841   0.8505
SVM      0.850  0.760    0.887   0.8774
XGB      0.889  0.785    0.904   0.9061
```

The max accuracy was by using Select From Model

9.8 Selected Feature

```
[ ]: all_inputs = df[df.columns[0:-1]].values
      sc = StandardScaler()
      all_inputs = sc.fit_transform(all_inputs)
      all_labels = df['class'].values
```

```

from sklearn.svm import LinearSVC
from sklearn.feature_selection import SelectFromModel

lsvc = LinearSVC(C=.09 , penalty="l1", dual=False).fit(all_inputs, all_labels)
model = SelectFromModel(lsvc, prefit=True)
all_inputs = model.transform(all_inputs)
print(all_inputs.shape)

from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳test_size=0.25, random_state=1)
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_4 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_4 = cross_val_score(model, X_train, y_train, cv=kfold,
    ↳scoring='accuracy')
    results_4.append(cv_results_4)
    names.append(name)
    print(f"{name}, {cv_results_4.mean()}, {cv_results_4.std()}")

```

```

(696, 14)
LR, 0.8716524216524217, 0.057349054491092656))
LDA, 0.8753561253561255, 0.05958861601904569))
KNN, 0.8657407407407408, 0.06845151790492156))
CART, 0.8467948717948717, 0.06182643886679986))
NB, 0.8504985754985755, 0.07195561951720005))
SVM, 0.865883190883191, 0.07546399654292588))
XGB, 0.9061253561253562, 0.05633700530500054))

```

```
[ ]: df.shape
```

```
[ ]: (696, 15)
```

10 5. Imbalanced Correction Methods

```

[ ]: all_inputs = df[df.columns[:-1]].values
    all_labels = df['class'].values
    pd.value_counts(df['class'])

```

```

[ ]: 0    456
    1    240

```

Name: class, dtype: int64

```
[ ]: all_inputs.shape
```


```
[ ]: (696, 14)
```

```
[ ]: from sklearn.utils.class_weight import compute_class_weight
class_weights = compute_class_weight('balanced', np.unique(all_labels),
    ↪all_labels)
print(class_weights)
```

```
[0.76315789 1.45      ]
```

```
[ ]: pd.value_counts(df['class']).plot(kind='pie', autopct='%.2f',fontsize=15,
    ↪colors= ['#66b3ff','#ff9999'])
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcadb549550>
```



10.1 1. SMOTE

```
[ ]: all_inputs = df[df.columns[:-1]].values
all_labels = df['class'].values
pd.value_counts(df['class'])
```

```
[ ]: 0    456
```

```
     1    240
```

Name: class, dtype: int64


```
[ ]: from imblearn.over_sampling import SMOTE
oversample = SMOTE()
all_inputs, all_labels = oversample.fit_resample(all_inputs, all_labels)
pd.value_counts(all_labels)

/usr/local/lib/python3.6/dist-packages/sklearn/externals/six.py:31:
FutureWarning: The module is deprecated in version 0.21 and will be removed in
version 0.23 since we've dropped support for Python 2.7. Please rely on the
official version of six (https://pypi.org/project/six/).
  "(https://pypi.org/project/six/).", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:144:
FutureWarning: The sklearn.neighbors.base module is deprecated in version 0.22
and will be removed in version 0.24. The corresponding classes / functions
should instead be imported from sklearn.neighbors. Anything that cannot be
imported from sklearn.neighbors is now part of the private API.
  warnings.warn(message, FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 1    456
      0    456
dtype: int64
```

```
[ ]: from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
↳test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_1 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_1 = cross_val_score(model, X_train, y_train, cv=kfold,
↳scoring='accuracy')
    results_1.append(cv_results_1)
    names.append(name)
    print(f"{name}, {cv_results_1.mean()}, {cv_results_1.std()}")
```

```
LR, 0.864327731092437, 0.07061647844343612))
LDA, 0.852563025210084, 0.07573870608521982))
KNN, 0.8803781512605042, 0.0689722359962584))
CART, 0.8847058823529412, 0.04496434416635351))
```

```
NB, 0.8264285714285714, 0.0796771941301821))
SVM, 0.8571428571428571, 0.08050885058768605))
XGB, 0.9139075630252099, 0.05052403517213921))
```

10.2 2. Border line SMOTE

```
[ ]: all_inputs = df[df.columns[:-1]].values
all_labels = df['class'].values
pd.value_counts(df['class'])
```

```
[ ]: 0    456
     1    240
     Name: class, dtype: int64
```

```
[ ]: from imblearn.over_sampling import BorderlineSMOTE
oversample = BorderlineSMOTE()
all_inputs, all_labels = oversample.fit_resample(all_inputs, all_labels)
pd.value_counts(all_labels)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 1    456
     0    456
     dtype: int64
```

```
[ ]: from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳ test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_2 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
```

```

cv_results_2 = cross_val_score(model, X_train, y_train, cv=kfold,
→scoring='accuracy')
results_2.append(cv_results_2)
names.append(name)
print(f"{name}, {cv_results_2.mean()}, {cv_results_2.std()}))")

```

```

LR, 0.8405882352941176, 0.06742518025605315))
LDA, 0.8333613445378152, 0.062351072476471266))
KNN, 0.8538655462184874, 0.07252003465854545))
CART, 0.8775630252100839, 0.0659260166576104))
NB, 0.8189495798319328, 0.06368865097025407))
SVM, 0.8744117647058822, 0.05982666568734233))
XGB, 0.9270588235294115, 0.03943281228312503))

```

10.3 3. SVM SMOTE

```

[: all_inputs = df[df.columns[:-1]].values
all_labels = df['class'].values
pd.value_counts(all_labels)

```

```

[: 0    456
   1    240
dtype: int64

```

```

[: from imblearn.over_sampling import SVMSMOTE
oversample = SVMSMOTE()
all_inputs, all_labels = oversample.fit_resample(all_inputs, all_labels)
pd.value_counts(all_labels)

```

```

/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)

```

```
warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 1    456
      0    456
      dtype: int64
```

```
[ ]: from sklearn.model_selection import train_test_split
      (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
      →test_size=0.25, random_state=1, stratify= all_labels )
      # Feature Scaling
      from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)

      results_3 = []; names = []; seed=42
      for name, model in models:
          kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
          cv_results_3 = cross_val_score(model, X_train, y_train, cv=kfold,
          →scoring='accuracy')
          results_3.append(cv_results_3)
          names.append(name)
          print(f"{name}, {cv_results_3.mean()}, {cv_results_3.std()}")
```

```
LR, 0.8568487394957984, 0.0632914683906705))
LDA, 0.8466386554621849, 0.052723642715918145))
KNN, 0.8670168067226891, 0.05101184696196572))
CART, 0.8876050420168067, 0.04973474942606156))
NB, 0.8190336134453782, 0.07572276088325611))
SVM, 0.8656302521008403, 0.06492808078200309))
XGB, 0.9227310924369746, 0.049563075789801826))
```

10.4 4. Adaptive Synthetic Sampling (ADASYN)

```
[ ]: all_inputs = df[df.columns[:-1]].values
      all_labels = df['class'].values
      pd.value_counts(all_labels)
```

```
[ ]: 0    456
      1    240
      dtype: int64
```

```
[ ]: from imblearn.over_sampling import ADASYN
      oversample = ADASYN(random_state=42)
      all_inputs, all_labels = oversample.fit_resample(all_inputs, all_labels)
```

```
pd.value_counts(all_labels)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 1    463
      0    456
      dtype: int64
```

```
[ ]: from sklearn.model_selection import train_test_split
      (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
      →test_size=0.25, random_state=1, stratify= all_labels )
      # Feature Scaling
      from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)

      results_4 = []; names = []; seed=42
      for name, model in models:
          kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
          cv_results_4 = cross_val_score(model, X_train, y_train, cv=kfold,
          →scoring='accuracy')
          results_4.append(cv_results_4)
          names.append(name)
          print(f"{name}, {cv_results_4.mean()}, {cv_results_4.std()}")
```

```
LR, 0.8057142857142857, 0.060236100684876004))
LDA, 0.7881512605042017, 0.05399177647682829))
KNN, 0.8735714285714286, 0.059343128313607926))
CART, 0.8809663865546218, 0.04492795274422015))
NB, 0.7763025210084035, 0.0543214753298728))
SVM, 0.8603361344537817, 0.05325013223997295))
XGB, 0.9173949579831933, 0.04475537476515333))
```

10.5 5. Random Over Sampler

```
[ ]: all_inputs = df[df.columns[:-1]].values
      all_labels = df['class'].values
      pd.value_counts(all_labels)
```

```
[ ]: 0    456
      1    240
      dtype: int64
```

```
[ ]: from imblearn.over_sampling import RandomOverSampler
oversample = RandomOverSampler(sampling_strategy='minority')
all_inputs, all_labels = oversample.fit_resample(all_inputs, all_labels)
pd.value_counts(all_labels)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 1    456
      0    456
      dtype: int64
```

```
[ ]: from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
↳test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_5 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_5 = cross_val_score(model, X_train, y_train, cv=kfold,
↳scoring='accuracy')
    results_5.append(cv_results_5)
    names.append(name)
    print(f"{name}, {cv_results_5.mean()}, {cv_results_5.std()}))")
```

```
LR, 0.8773949579831932, 0.05463988066938839))
LDA, 0.8729831932773108, 0.056153241836470545))
KNN, 0.8641596638655462, 0.06511888006068071))
CART, 0.9124369747899159, 0.04493213739132524))
NB, 0.8277731092436975, 0.06759065188336488))
SVM, 0.8788655462184873, 0.05422237569869624))
XGB, 0.9372268907563024, 0.05165921533119705))
```

10.6 6. Random Under Sampler

```
[ ]: all_inputs = df[df.columns[:-1]].values
all_labels = df['class'].values
pd.value_counts(all_labels)
```

```
[ ]: 0    456
      1    240
```

dtype: int64

```
[ ]: from imblearn.under_sampling import RandomUnderSampler
undersample = RandomUnderSampler()
all_inputs, all_labels = undersample.fit_resample(all_inputs, all_labels)
pd.value_counts(all_labels)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 1    240
      0    240
dtype: int64
```

```
[ ]: from sklearn.model_selection import train_test_split
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
↳test_size=0.25, random_state=1, stratify= all_labels )
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler( )
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

results_6 = []; names = []; seed=42
for name, model in models:
    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    cv_results_6 = cross_val_score(model, X_train, y_train, cv=kfold,
↳scoring='accuracy')
    results_6.append(cv_results_6)
    names.append(name)
    print(f"{name}, {cv_results_6.mean()}, {cv_results_6.std()}")
```

```
LR, 0.8388888888888889, 0.08940820521906158))
LDA, 0.8388888888888889, 0.09606453592105879))
KNN, 0.8194444444444444, 0.08216777476527243))
CART, 0.8416666666666668, 0.09005313931915183))
NB, 0.8111111111111111, 0.09196080754026026))
SVM, 0.8472222222222223, 0.07632573146685602))
XGB, 0.8694444444444445, 0.07301910793386061))
```

10.7 7. Combining Random Oversampling and Undersampling

```
[ ]: all_inputs = df[df.columns[:-1]].values
all_labels = df['class'].values
pd.value_counts(all_labels)
```

```
[ ]: 0    456
      1    240
      dtype: int64
```

```
[ ]: under = RandomUnderSampler(sampling_strategy=.7)
      all_inputs, all_labels = under.fit_resample(all_inputs, all_labels)
      pd.value_counts(all_labels)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 0    342
      1    240
      dtype: int64
```

```
[ ]: over = RandomOverSampler()
      all_inputs, all_labels = over.fit_resample(all_inputs, all_labels)
      pd.value_counts(all_labels)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 1    342
      0    342
      dtype: int64
```

```
[ ]: from sklearn.model_selection import train_test_split
      (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
      ↪test_size=0.25, random_state=1, stratify= all_labels )
      # Feature Scaling
      from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)

      results_7 = []; names = []; seed=42
      for name, model in models:
          kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
          cv_results_7 = cross_val_score(model, X_train, y_train, cv=kfold,
          ↪scoring='accuracy')
          results_7.append(cv_results_7)
          names.append(name)
          print(f"{name}, {cv_results_7.mean()}, {cv_results_7.std()}")
```



```

LR, 0.8576923076923076, 0.07534300264349224))
LDA, 0.8537692307692308, 0.07913508037997789))
KNN, 0.8261538461538462, 0.07429670248402684))
CART, 0.8772307692307691, 0.06310055842362079))
NB, 0.8206153846153846, 0.08142139637086955))
SVM, 0.8516923076923077, 0.07258604361590905))
XGB, 0.8944615384615385, 0.07443896669966349))

```

10.8 Classification Comparison of Feature Selection Methods

```
[ ]: classification_comparison= pd.DataFrame(index=[i for i in names])
```

```
[ ]: dfresults = pd.DataFrame(results_1)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy']= df_mean
classification_comparison['sd']= df_sd
```

```
[ ]: dfresults = pd.DataFrame(results_2)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_2']= df_mean
classification_comparison['sd_2']= df_sd
```

```
[ ]: dfresults = pd.DataFrame(results_3)
dfresults=dfresults.T
```

```

dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_3']= df_mean
classification_comparison['sd_3']= df_sd

```

```

[ ]: dfresults = pd.DataFrame(results_4)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_4']= df_mean
classification_comparison['sd_4']= df_sd

```

```

[ ]: dfresults = pd.DataFrame(results_5)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_5']= df_mean
classification_comparison['sd_5']= df_sd

```

```
[ ]: dfresults = pd.DataFrame(results_6)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_6']= df_mean
classification_comparison['sd_6']= df_sd
```

```
[ ]: dfresults = pd.DataFrame(results_7)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy_7']= df_mean
classification_comparison['sd_7']= df_sd
```

```
[ ]: values= ['Acc', 'SD']
```

```
[ ]: FS_Methods=['1_SMOTE', '2_BL-SMOTE', '3_SVM-SMOTE',
                '4_ADASYN', '5_ROS', '6_RUS', '7_CR(O-U)S']
```

```
[ ]: idx = pd.MultiIndex.from_product([FS_Methods, values],
                                     names=['FS_Methods', 'values'])

classification_comparison.columns = idx
classification_comparison
```


```
[ ]: FS_Methods 1_SMOTE          2_BL-SMOTE  ...   6_RUS 7_CR(O-U)S
values      Acc      SD          Acc  ...      SD      Acc      SD
LR          0.8643  0.0725    0.8406  ...  0.0917    0.8577  0.0773
LDA          0.8526  0.0777    0.8334  ...  0.0986    0.8538  0.0812
```

KNN	0.8804	0.0708	0.8539	...	0.0843	0.8262	0.0762
CART	0.8847	0.0461	0.8776	...	0.0924	0.8772	0.0647
NB	0.8264	0.0817	0.8189	...	0.0943	0.8206	0.0835
SVM	0.8571	0.0826	0.8744	...	0.0783	0.8517	0.0745
XGB	0.9139	0.0518	0.9271	...	0.0749	0.8945	0.0764

[7 rows x 14 columns]

```
[ ]: classification_comparison.to_csv('Classification Comparison of Feature_
    ↳Selection Methods.csv', index= False)

[ ]: plt.figure(figsize=(10,7))
    ax = plt.subplot(211)
    classification_comparison.xs('Acc', axis=1, level='values').plot( ax=ax)
    plt.title("Accuracy")
    plt.legend(shadow=True, frameon=True, fancybox=True, title='FS_Methods',
    ↳bbox_to_anchor=(1.05, 1), loc='upper left')
    ax = plt.subplot(212)
    classification_comparison.xs('SD', axis=1, level='values').plot( ax=ax)
    plt.title("Standard Deviation")
    plt.xlabel("Classification Algorithms")
    plt.legend(shadow=True, frameon=True, fancybox=True, title='FS_Methods',
    ↳bbox_to_anchor=(1.05, 1), loc='upper left')
    plt.tight_layout()
```



```
[ ]: classification_comparison.T.max()
```

```
[ ]: LR      0.8774  
     LDA      0.8730  
     KNN      0.8804  
     CART      0.9124  
     NB       0.8278  
     SVM      0.8789  
     XGB      0.9372  
     dtype: float64
```

```
[ ]: Imbalance= classification_comparison.T.max()
```

```
[ ]: classification_comparison.T['CART']
```

```
[ ]: FS_Methods  values  
     1_SMOTE     Acc      0.8847  
           SD      0.0461  
     2_BL-SMOTE  Acc      0.8776  
           SD      0.0676  
     3_SVM-SMOTE Acc      0.8876  
           SD      0.0510  
     4_ADASYN    Acc      0.8810  
           SD      0.0461  
     5_ROS       Acc      0.9124  
           SD      0.0461  
     6_RUS       Acc      0.8417  
           SD      0.0924  
     7_CR(0-U)S  Acc      0.8772  
           SD      0.0647  
     Name: CART, dtype: float64
```

```
[ ]: classification_comparison.T['XGB']
```

```
[ ]: FS_Methods  values  
     1_SMOTE     Acc      0.9139  
           SD      0.0518  
     2_BL-SMOTE  Acc      0.9271  
           SD      0.0405  
     3_SVM-SMOTE Acc      0.9227  
           SD      0.0509  
     4_ADASYN    Acc      0.9174  
           SD      0.0459  
     5_ROS       Acc      0.9372  
           SD      0.0530  
     6_RUS       Acc      0.8694  
           SD      0.0749  
     7_CR(0-U)S  Acc      0.8945  
           SD      0.0764
```

Name: XGB, dtype: float64

```
[ ]: methods_comparison['Imbalance']= Imbalance
      methods_comparison
```

```
[ ]:      Fill_NaN  Add_F  Outlier  Feature  Imbalance
      LR          0.786  0.765    0.891   0.8735    0.8774
      LDA          0.790  0.767    0.895   0.8754    0.8730
      KNN          0.817  0.757    0.890   0.8657    0.8804
      CART         0.880  0.724    0.888   0.8619    0.9124
      NB           0.780  0.746    0.841   0.8505    0.8278
      SVM          0.850  0.760    0.887   0.8774    0.8789
      XGB          0.889  0.785    0.904   0.9061    0.9372
```

10.9 Selected imbalance methods

The max accuracy was by using Random Over Sampler

```
[ ]: all_inputs = df[df.columns[:-1]].values
      all_labels = df['class'].values
      all_inputs.shape
```

```
[ ]: (696, 14)
```

```
[ ]: from imblearn.over_sampling import BorderlineSMOTE
      oversample = BorderlineSMOTE()
      all_inputs, all_labels = oversample.fit_resample(all_inputs, all_labels)
      pd.value_counts(all_labels)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
```

```
[ ]: 1    456
      0    456
      dtype: int64
```

11 Step 5: Building the classifier

11.1 1. XGBoost Classifier

11.1.1 Hyperparameter Randomized Search CV

```
[ ]: from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import StratifiedKFold

params = {
    'min_child_weight': [1, 5, 10],
    'gamma': [0.5, 1, 1.5, 2, 5],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'max_depth': [3, 4, 5]
}

xgb = XGBClassifier(learning_rate=0.02, n_estimators=600, objective='binary:
    ↳logistic',
                    silent=True, nthread=1)

folds = 3
param_comb = 5

skf = StratifiedKFold(n_splits=folds, shuffle = True, random_state = 1001)

random_search = RandomizedSearchCV(xgb, param_distributions=params,
                                   n_iter=param_comb, scoring='roc_auc',
                                   n_jobs=4, cv=skf.
    ↳split(all_inputs,all_labels),
                                   verbose=3, random_state=1001 )
random_search.fit(all_inputs,all_labels)
```

Fitting 3 folds for each of 5 candidates, totalling 15 fits

[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=4)]: Done 15 out of 15 | elapsed: 6.6s finished

```
[ ]: RandomizedSearchCV(cv=<generator object _BaseKFold.split at 0x7fcad356cba0>,
    error_score=nan,
    estimator=XGBClassifier(base_score=0.5, booster='gbtree',
        colsample_bylevel=1,
        colsample_bynode=1,
        colsample_bytree=1, gamma=0,
        learning_rate=0.02, max_delta_step=0,
        max_depth=3, min_child_weight=1,
        missing=None, n_estimators=600,
        n_jobs=1, nthread=1,
```

```

        objective='binary:logist...
        reg_lambda=1, scale_pos_weight=1,
        seed=None, silent=True, subsample=1,
        verbosity=1),
    iid='deprecated', n_iter=5, n_jobs=4,
    param_distributions={'colsample_bytree': [0.6, 0.8, 1.0],
                        'gamma': [0.5, 1, 1.5, 2, 5],
                        'max_depth': [3, 4, 5],
                        'min_child_weight': [1, 5, 10],
                        'subsample': [0.6, 0.8, 1.0]},
    pre_dispatch='2*n_jobs', random_state=1001, refit=True,
    return_train_score=False, scoring='roc_auc', verbose=3)

```

```

[:]: print('\n Best estimator:')
print(random_search.best_estimator_)
print('\n Best normalized gini score for %d-fold search with %d parameter_
      ↪combinations:' % (folds, param_comb))
print(random_search.best_score_ * 2 - 1)
print('\n Best hyperparameters:')
print(random_search.best_params_)
results = pd.DataFrame(random_search.cv_results_)
results.to_csv('xgb-random-grid-search-results-01.csv', index=False)

```

Best estimator:

```

XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=0.8, gamma=1.5,
              learning_rate=0.02, max_delta_step=0, max_depth=5,
              min_child_weight=1, missing=None, n_estimators=600, n_jobs=1,
              nthread=1, objective='binary:logistic', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=True, subsample=0.6, verbosity=1)

```

Best normalized gini score for 3-fold search with 5 parameter combinations:
0.92321675900277

Best hyperparameters:

```

{'subsample': 0.6, 'min_child_weight': 1, 'max_depth': 5, 'gamma': 1.5,
 'colsample_bytree': 0.8}

```

11.1.2 Confusion Matrix & Tuning

```

[:]: (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
      ↪test_size=0.25, random_state=1)
from sklearn.preprocessing import RobustScaler
sc = RobustScaler()
X_train = sc.fit_transform(X_train)

```



```

X_test = sc.transform(X_test)

cls = XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=0.8, gamma=1.5,
                    learning_rate=0.02, max_delta_step=0, max_depth=5,
                    min_child_weight=1, missing=None, n_estimators=600,
                    n_jobs=1,
                    nthread=1, objective='binary:logistic', random_state=0,
                    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                    silent=True, subsample=0.6, verbosity=1)
cls.fit(X_train,y_train)
y_pred = cls.predict(X_test)
print("----- Accuracy -----\\n")
print(accuracy_score(y_test,y_pred))
print('-----')
print("----- Confusion Matrix -----\\n")
print(confusion_matrix(y_test,y_pred))
print("----- Classification Report----- \\n")
print(classification_report(y_test,y_pred))
print('-----')

```

----- Accuracy -----

0.9342105263157895

----- Confusion Matrix -----

```
[[100   6]
 [  9 113]]
```

----- Classification Report-----

	precision	recall	f1-score	support
0	0.92	0.94	0.93	106
1	0.95	0.93	0.94	122
accuracy			0.93	228
macro avg	0.93	0.93	0.93	228
weighted avg	0.93	0.93	0.93	228

```

[ ]: XGB=[]
     XGB.append(0.934)

```

11.1.3 Preprocessing Methods Comparison

1. Standard Scaler
2. Min Max Scaler
3. Max Abs Scaler
4. Robust Scaler
5. Power Transformer
6. Quantile Transformer (uniform output)
7. Quantile Transformer (Gaussian output)
8. Normalizer

```
[ ]: from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import MaxAbsScaler
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import PowerTransformer
from sklearn.preprocessing import QuantileTransformer
from sklearn.preprocessing import Normalizer
scalers = []
scalers.append(('1.SS', StandardScaler()))
scalers.append(('2.MMS', MinMaxScaler()))
scalers.append(('3.MAS', MaxAbsScaler()))
scalers.append(('4.RPS', RobustScaler()))
scalers.append(('5.PT', PowerTransformer()))
scalers.append(('6.QTG', QuantileTransformer()))
scalers.append(('7.QTN', QuantileTransformer(output_distribution='normal')))
scalers.append(('8.NRM', Normalizer()))
```

```
[ ]: all_inputs = df[df.columns[:-1]].values
all_labels = df['class'].values

results_1 = []; names = []; seed=42
for name, scaler in scalers:

    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)

    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    model = cls
    cv_results_1 = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_1.append(cv_results_1)
    names.append(name)
    print(f"{name}, {cv_results_1.mean()}, {cv_results_1.std()}")
```

```
1.SS, 0.925546218487395, 0.04157617695531443))
2.MMS, 0.927016806722689, 0.040529364721187754))
3.MAS, 0.927016806722689, 0.040529364721187754))
```

```
4.RPS, 0.925546218487395, 0.04157617695531443))
5.PT, 0.927016806722689, 0.040529364721187754))
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/preprocessing/_data.py:2357:
UserWarning: n_quantiles (1000) is greater than the total number of samples
(684). n_quantiles is set to n_samples.
% (self.n_quantiles, n_samples))
```

```
6.QTG, 0.9299579831932773, 0.039297081894647666))
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/preprocessing/_data.py:2357:
UserWarning: n_quantiles (1000) is greater than the total number of samples
(684). n_quantiles is set to n_samples.
% (self.n_quantiles, n_samples))
```

```
7.QTN, 0.9299579831932773, 0.039297081894647666))
```

```
8.NRM, 0.9197478991596638, 0.05022616738997399))
```

```
[ ]: classification_comparison= pd.DataFrame(index=[i for i in names])
dfresults = pd.DataFrame(results_1)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy']= df_mean
classification_comparison['sd']= df_sd

classification_comparison
```


```
[ ]:      Accuracy      sd
1.SS      0.9255  0.0427
2.MMS      0.9270  0.0416
3.MAS      0.9270  0.0416
4.RPS      0.9255  0.0427
5.PT      0.9270  0.0416
6.QTG      0.9300  0.0403
7.QTN      0.9300  0.0403
8.NRM      0.9197  0.0515
```

11.1.4 Cross-validation

```
[ ]: from sklearn.model_selection import cross_val_score

# cross_val_score returns a list of the scores, which we can visualize
# to get a reasonable estimate of our classifier's performance
cv_scores = cross_val_score(cls, all_inputs, all_labels, cv=15)
plt.hist(cv_scores)
plt.title('Average score: {}'.format(np.mean(cv_scores)))
;
```

```
[ ]: ''
```



```
[ ]: XGB.append(np.mean(cv_scores).round(3))
XGB
```

```
[ ]: [0.934, 0.904]
```

11.2 2. Extra Trees Classifier

```
[ ]: from sklearn.ensemble import ExtraTreesClassifier

[ ]: from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import GridSearchCV

ETC = ExtraTreesClassifier(max_features=13,n_estimators= 520)
```

```

parameter_grid = {

    'max_leaf_nodes' : [5000,700]}

cross_validation = StratifiedKFold(n_splits=10)

grid_search = GridSearchCV(ETC,
                            param_grid=parameter_grid,
                            cv=cross_validation, n_jobs=4, )

grid_search.fit(all_inputs, all_labels)
print('Best score: {}'.format(grid_search.best_score_))
print('Best parameters: {}'.format(grid_search.best_params_))

```

Best score: 0.8907660455486542
 Best parameters: {'max_leaf_nodes': 700}

11.2.1 Hyperparameter Randomized Search CV

11.2.2 Confusion Matrix

```

[ ]: (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳test_size=0.25, random_state=1)
from sklearn.preprocessing import RobustScaler
sc = RobustScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

cls = ExtraTreesClassifier(n_estimators=9, max_features=13,
                           criterion= 'entropy',
                           min_samples_split=3,
                           max_samples= .5, n_jobs=100)

cls.fit(X_train,y_train)
y_pred = cls.predict(X_test)
print("----- Accuracy ----- \n")
print(accuracy_score(y_test,y_pred))
print('-----')
print("----- Confusion Matrix ----- \n")
print(confusion_matrix(y_test,y_pred))
print("----- Classification Report----- \n")
print(classification_report(y_test,y_pred))
print('-----')

```

----- Accuracy -----

0.8735632183908046

----- Confusion Matrix -----

```
[[101  9]
 [ 13 51]]
```

----- Classification Report-----

	precision	recall	f1-score	support
0	0.89	0.92	0.90	110
1	0.85	0.80	0.82	64
accuracy			0.87	174
macro avg	0.87	0.86	0.86	174
weighted avg	0.87	0.87	0.87	174

```
[ ]: ETsC=[]
     ETsC.append(0.873)
```

11.2.3 Preprocessing Methods Comparison

1. Standard Scaler
2. Min Max Scaler
3. Max Abs Scaler
4. Robust Scaler
5. Power Transformer
6. Quantile Transformer (uniform output)
7. Quantile Transformer (Gaussian output)
8. Normalizer

```
[ ]: from sklearn.preprocessing import StandardScaler
     from sklearn.preprocessing import MinMaxScaler
     from sklearn.preprocessing import MaxAbsScaler
     from sklearn.preprocessing import RobustScaler
     from sklearn.preprocessing import PowerTransformer
     from sklearn.preprocessing import QuantileTransformer
     from sklearn.preprocessing import Normalizer
     scalers = []
     scalers.append(('1.SS', StandardScaler()))
     scalers.append(('2.MMS', MinMaxScaler()))
     scalers.append(('3.MAS', MaxAbsScaler()))
     scalers.append(('4.RPS', RobustScaler()))
     scalers.append(('5.PT', PowerTransformer()))
     scalers.append(('6.QTG', QuantileTransformer()))
     scalers.append(('7.QTN', QuantileTransformer(output_distribution='normal')))
     scalers.append(('8.NRM', Normalizer()))
```

```
[ ]: all_inputs = df[df.columns[:-1]].values
all_labels = df['class'].values

results_1 = []; names = []; seed=42
for name, scaler in scalers:

    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)

    kfold = KFold(n_splits=20, random_state=seed, shuffle=True)
    model = cls
    cv_results_1 = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_1.append(cv_results_1)
    names.append(name)
    print(f"{name}, {cv_results_1.mean()}, {cv_results_1.std()}")"
```

```
1.SS, 0.8562678062678064, 0.06428574649529842))
2.MMS, 0.8792735042735043, 0.057460359946266484))
3.MAS, 0.8792735042735043, 0.06472454657741178))
4.RPS, 0.8907407407407406, 0.06121230307312289))
5.PT, 0.8619658119658121, 0.061471651767881295))
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/preprocessing/_data.py:2357:
UserWarning: n_quantiles (1000) is greater than the total number of samples
(522). n_quantiles is set to n_samples.
% (self.n_quantiles, n_samples))
```

```
6.QTG, 0.8735042735042736, 0.07213746100838564))
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/preprocessing/_data.py:2357:
UserWarning: n_quantiles (1000) is greater than the total number of samples
(522). n_quantiles is set to n_samples.
% (self.n_quantiles, n_samples))
```

```
7.QTN, 0.8698005698005697, 0.07904943336596147))
8.NRM, 0.8619658119658119, 0.06383276556156207))
```

```
[ ]: classification_comparison= pd.DataFrame(index=[i for i in names])
dfresults = pd.DataFrame(results_1)
dfresults=dfresults.T
dfresults.columns=names
df_mean=[]
df_sd=[]
for i in dfresults.columns:
    d= dfresults[i].mean()
    df_mean.append(d)
df_mean= [round(num, 4) for num in df_mean]
```

```

for i in dfresults.columns:
    n= dfresults[i].std()
    df_sd.append(n)
df_sd= [round(num, 4) for num in df_sd]

classification_comparison['Accuracy']= df_mean
classification_comparison['sd']= df_sd

classification_comparison

```

```

[:]:      Accuracy      sd
1.SS      0.8563  0.0660
2.MMS      0.8793  0.0590
3.MAS      0.8793  0.0664
4.RPS      0.8907  0.0628
5.PT       0.8620  0.0631
6.QTG      0.8735  0.0740
7.QTN      0.8698  0.0811
8.NRM      0.8620  0.0655

```

11.2.4 Cross-validation

```

[:]: from sklearn.model_selection import cross_val_score


# cross_val_score returns a list of the scores, which we can visualize
# to get a reasonable estimate of our classifier's performance
cv_scores = cross_val_score(cls, all_inputs, all_labels, cv=15)
plt.hist(cv_scores)
plt.title('Average score: {}'.format(np.mean(cv_scores)))
;

```

```

[:]: ''

```

```
[ ]: ETsC.append(np.mean(cv_scores).round(3))  
ETsC
```

```
[ ]: [0.873, 0.888]
```

11.3 3. LGBM Classifier

11.3.1 Confusion Matrix & Tuning

```
[ ]: from lightgbm import LGBMClassifier  
[ ]: from sklearn.model_selection import RandomizedSearchCV  
import lightgbm as lgb  
  
rs_params = {  
    'bagging_fraction': (0.5, 0.8),  
    'bagging_frequency': (5, 8),  
  
    'feature_fraction': (0.5, 0.8),  
    'max_depth': (10, 13),  
    'min_data_in_leaf': (90, 120),  
    'num_leaves': (1200, 1550)  
}
```

```

lgb = LGBMClassifier()

folds = 3
param_comb = 5

skf = StratifiedKFold(n_splits=folds, shuffle = True, random_state = 1001)

random_search = RandomizedSearchCV(lgb, param_distributions=params,
                                   n_iter=param_comb, scoring='roc_auc',
                                   n_jobs=4, cv=skf.
    →split(all_inputs,all_labels),
                                   verbose=3, random_state=1001 )
random_search.fit(all_inputs,all_labels)

```

Fitting 3 folds for each of 5 candidates, totalling 15 fits

[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=4)]: Done 15 out of 15 | elapsed: 2.9s finished

```

[: RandomizedSearchCV(cv=<generator object _BaseKFold.split at 0x7fcad3744678>,
                     error_score=nan,
                     estimator=LGBMClassifier(boosting_type='gbdt',
                                                class_weight=None,
                                                colsample_bytree=1.0,
                                                importance_type='split',
                                                learning_rate=0.1, max_depth=-1,
                                                min_child_samples=20,
                                                min_child_weight=0.001,
                                                min_split_gain=0.0,
                                                n_estimators=100, n_jobs=-1,
                                                num_leaves=31, objective=None,
                                                ran...
                                                subsample=1.0,
                                                subsample_for_bin=200000,
                                                subsample_freq=0),
                     iid='deprecated', n_iter=5, n_jobs=4,
                     param_distributions={'ccp_alpha': [0.1, 0.5, 0.9],
                                           'colsample_bytree': [0.6, 0.8, 0.1],
                                           'max_depth': [3, 4, 5],
                                           'min_samples_split': [2, 3, 5, 8],
                                           'verbose': [0.5, 1, 1.5, 2, 5]},
                     pre_dispatch='2*n_jobs', random_state=1001, refit=True,
                     return_train_score=False, scoring='roc_auc', verbose=3)

```

```

[: print('\n Best estimator:')
   print(random_search.best_estimator_)
   print('\n Best normalized gini score for %d-fold search with %d parameter_
    →combinations:' % (folds, param_comb))

```

```

print(random_search.best_score_ * 2 - 1)
print('\n Best hyperparameters:')
print(random_search.best_params_)
results = pd.DataFrame(random_search.cv_results_)
results.to_csv('xgb-random-grid-search-results-01.csv', index=False)

```

Best estimator:

```

LGBMClassifier(boosting_type='gbdt', ccp_alpha=0.5, class_weight=None,
               colsample_bytree=0.6, importance_type='split', learning_rate=0.1,
               max_depth=3, min_child_samples=20, min_child_weight=0.001,
               min_samples_split=3, min_split_gain=0.0, n_estimators=100,
               n_jobs=-1, num_leaves=31, objective=None, random_state=None,
               reg_alpha=0.0, reg_lambda=0.0, silent=True, subsample=1.0,
               subsample_for_bin=200000, subsample_freq=0, verbose=5)

```

Best normalized gini score for 3-fold search with 5 parameter combinations:
0.9081688596491226

Best hyperparameters:

```

{'verbose': 5, 'min_samples_split': 3, 'max_depth': 3, 'colsample_bytree': 0.6,
 'ccp_alpha': 0.5}

```

```

[: (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳ test_size=0.25, random_state=1)
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

cls = LGBMClassifier(num_leaves= 55,
                    learning_rate = .3,
                    subsample_for_bin = 233,
                    importance_type = 'gain')
cls.fit(X_train,y_train)
y_pred = cls.predict(X_test)
print("----- Accuracy ----- \n")
print(accuracy_score(y_test,y_pred))
print('-----')
print("----- Confusion Matrix ----- \n")
print(confusion_matrix(y_test,y_pred))
print("----- Classification Report----- \n")
print(classification_report(y_test,y_pred))
print('-----')

```

----- Accuracy -----

0.9022988505747126

----- Confusion Matrix -----

```
[[102   8]
 [  9  55]]
```

----- Classification Report -----

	precision	recall	f1-score	support
0	0.92	0.93	0.92	110
1	0.87	0.86	0.87	64
accuracy			0.90	174
macro avg	0.90	0.89	0.89	174
weighted avg	0.90	0.90	0.90	174


```
[ ]: LGBM=[]
      LGBM.append(0.902)
```

11.3.2 Cross-validation

```
[ ]: from sklearn.model_selection import cross_val_score

      # cross_val_score returns a list of the scores, which we can visualize
      # to get a reasonable estimate of our classifier's performance
      cv_scores = cross_val_score(cls, all_inputs, all_labels, cv=15)
      plt.hist(cv_scores)
      plt.title('Average score: {}'.format(np.mean(cv_scores)))
      ;
```

```
[ ]: ''
```



```
[ ]: LGBM.append(np.mean(cv_scores).round(3))
      LGBM
```

```
[ ]: [0.902, 0.894]
```

11.4 4. Decision Tree Classifier

11.4.1 Confusion Matrix & Tuning

```
[ ]: from sklearn.model_selection import StratifiedKFold
      from sklearn.model_selection import GridSearchCV

      decision_tree_classifier = DecisionTreeClassifier()

      parameter_grid = {'criterion': ['gini', 'entropy'],
                        'splitter': ['best', 'random'],
                        'max_depth': list(range(21,55,2)),
                        'max_features': list(range(1,14))}

      cross_validation = StratifiedKFold(n_splits=10)

      grid_search = GridSearchCV(decision_tree_classifier,
                                param_grid=parameter_grid,
                                cv=cross_validation)
```

```

grid_search.fit(all_inputs, all_labels)
print('Best score: {}'.format(grid_search.best_score_))
print('Best parameters: {}'.format(grid_search.best_params_))

```

Best score: 0.8863975155279503
 Best parameters: {'criterion': 'gini', 'max_depth': 47, 'max_features': 11, 'splitter': 'best'}

```

[ ]: (X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
    ↳ test_size=0.25, random_state=1)
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

cls = DecisionTreeClassifier(criterion='entropy',
                             max_depth=21,
                             max_features=11,
                             splitter='best')

cls.fit(X_train,y_train)
y_pred = cls.predict(X_test)
print("----- Accuracy ----- \n")
print(accuracy_score(y_test,y_pred))
print('-----')
print("----- Confusion Matrix ----- \n")
print(confusion_matrix(y_test,y_pred))
print("----- Classification Report----- \n")
print(classification_report(y_test,y_pred))
print('-----')

```

----- Accuracy -----

0.9022988505747126

----- Confusion Matrix -----

```

[[102   8]
 [  9  55]]

```

----- Classification Report-----

	precision	recall	f1-score	support
0	0.92	0.93	0.92	110
1	0.87	0.86	0.87	64
accuracy			0.90	174
macro avg	0.90	0.89	0.89	174


weighted avg	0.90	0.90	0.90	174
--------------	------	------	------	-----

```
[ ]: DTC=[]  
DTC.append(0.902)
```

11.4.2 Cross-validation

```
[ ]: from sklearn.model_selection import cross_val_score  
  
# cross_val_score returns a list of the scores, which we can visualize  
# to get a reasonable estimate of our classifier's performance  
cv_scores = cross_val_score(cls, all_inputs, all_labels, cv=15)  
plt.hist(cv_scores, bins=9, stacked=True)  
plt.title('Average score: {}'.format(np.mean(cv_scores)))  
;
```

```
[ ]: ''
```



```
[ ]: DTC.append(np.mean(cv_scores).round(3))  
DTC
```

```
[ ]: [0.902, 0.875]
```

11.5 Comparison

```
[ ]: compile= pd.DataFrame(ETsC, index=['Acc','Val_acc'], columns=['ETsC'])
compile['XGB']= XGB
compile['LGBM']= LGBM
compile['DTC']= DTC
compile= compile.T
compile
```

```
[ ]:      Acc  Val_acc
ETsC  0.873    0.888
XGB    0.934    0.904
LGBM   0.902    0.894
DTC    0.902    0.875
```

12 Selected model


```
[ ]: cls = XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=0.8, gamma=1.5,
                        learning_rate=0.02, max_delta_step=0, max_depth=5,
                        min_child_weight=1, missing=None, n_estimators=600,
                        n_jobs=1,
                        nthread=1, objective='binary:logistic', random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                        silent=True, subsample=0.6, verbosity=1)
model_accuracies = []

for repetition in range(1000):
    (training_inputs,
     testing_inputs,
     training_classes,
     testing_classes) = train_test_split(all_inputs, all_labels, test_size=0.30)

    cls.fit(training_inputs, training_classes)
    classifier_accuracy = cls.score(testing_inputs, testing_classes)
    model_accuracies.append(classifier_accuracy)

plt.hist(model_accuracies)
plt.title('Average score: {}'.format(np.mean(classifier_accuracy)))
;
```

```
[ ]: ''
```

```
[ ]: dt_scores = cross_val_score(cls, all_inputs, all_labels, cv=10)

sb.boxplot(dt_scores, color='gray', )
sb.stripplot(dt_scores, jitter=True, color='black')
;
```


/usr/local/lib/python3.6/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

/usr/local/lib/python3.6/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

```
[ ]: ''
```




```
[ ]: sb.distplot(dt_scores)
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in a  
future version. Please adapt your code to use either `displot` (a figure-level  
function with similar flexibility) or `histplot` (an axes-level function for  
histograms).
```

```
warnings.warn(msg, FutureWarning)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcadb5abd30>
```

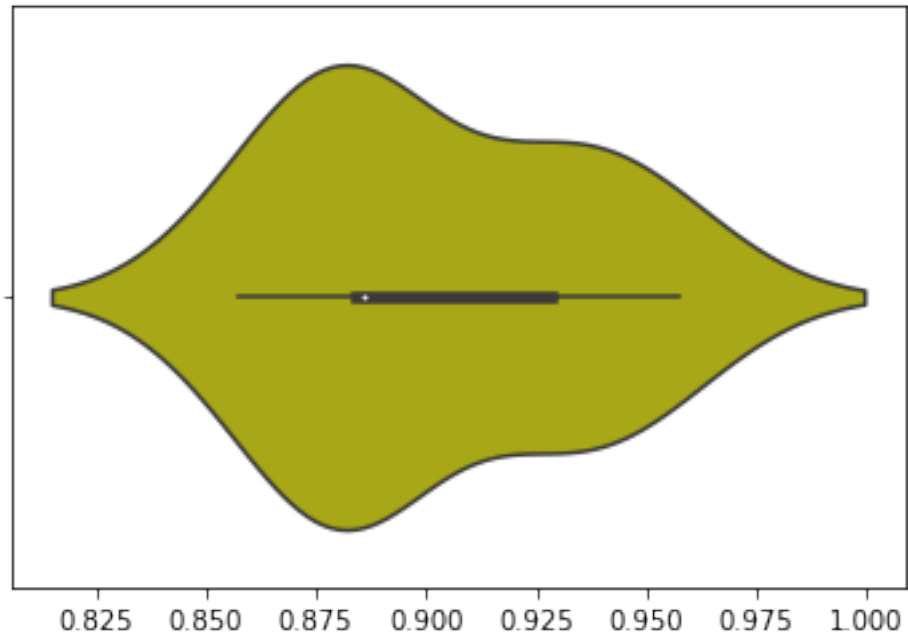


```
[ ]: sb.violinplot(dt_scores, color='y')
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/_decorators.py:43: FutureWarning:  
Pass the following variable as a keyword arg: x. From version 0.12, the only  
valid positional argument will be `data`, and passing other arguments without an  
explicit keyword will result in an error or misinterpretation.
```

```
FutureWarning
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcadd11b400>
```



13 Using Neural Networks

```
[ ]: import tensorflow as tf
(X_train, X_test, y_train, y_test) = train_test_split(all_inputs, all_labels,
                                                    test_size=0.2,
                                                    random_state=1,
                                                    stratify= all_labels )

# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

[ ]: def build_model():
    model = tf.keras.Sequential([
        tf.keras.layers.Dense(14, activation='relu', input_shape=[14]),
        tf.keras.layers.Dense(4, activation='relu'),
        tf.keras.layers.Dense(1, activation='sigmoid')
    ])

    optimizer = tf.keras.optimizers.RMSprop(0.01)

    model.compile(loss='binary_crossentropy', optimizer=optimizer,
    metrics=['accuracy'])
    return model
```

```
model = build_model()
```

```
[ ]: del model
```

```
[ ]: Epochs = 75  
history = model.fit(X_train, y_train, epochs=Epochs,  
                    validation_split=0.2,)
```

Epoch 1/75

14/14 [=====] - 1s 15ms/step - loss: 0.7363 - accuracy:
0.5698 - val_loss: 0.6165 - val_accuracy: 0.8036

Epoch 2/75

14/14 [=====] - 0s 4ms/step - loss: 0.5889 - accuracy:
0.8301 - val_loss: 0.5116 - val_accuracy: 0.8571

Epoch 3/75

14/14 [=====] - 0s 4ms/step - loss: 0.4434 - accuracy:
0.8499 - val_loss: 0.4022 - val_accuracy: 0.8661

Epoch 4/75

14/14 [=====] - 0s 4ms/step - loss: 0.3221 - accuracy:
0.8853 - val_loss: 0.3796 - val_accuracy: 0.8839

Epoch 5/75

14/14 [=====] - 0s 4ms/step - loss: 0.3168 - accuracy:
0.8904 - val_loss: 0.3835 - val_accuracy: 0.8571

Epoch 6/75

14/14 [=====] - 0s 4ms/step - loss: 0.2661 - accuracy:
0.8838 - val_loss: 0.3870 - val_accuracy: 0.8571

Epoch 7/75

14/14 [=====] - 0s 4ms/step - loss: 0.2607 - accuracy:
0.8998 - val_loss: 0.4030 - val_accuracy: 0.8661

Epoch 8/75

14/14 [=====] - 0s 4ms/step - loss: 0.2512 - accuracy:
0.9153 - val_loss: 0.4117 - val_accuracy: 0.8750

Epoch 9/75

14/14 [=====] - 0s 4ms/step - loss: 0.2173 - accuracy:
0.9236 - val_loss: 0.4176 - val_accuracy: 0.8571

Epoch 10/75

14/14 [=====] - 0s 4ms/step - loss: 0.2542 - accuracy:
0.8982 - val_loss: 0.4247 - val_accuracy: 0.8482

Epoch 11/75

14/14 [=====] - 0s 4ms/step - loss: 0.2214 - accuracy:
0.9120 - val_loss: 0.4622 - val_accuracy: 0.8482

Epoch 12/75

14/14 [=====] - 0s 4ms/step - loss: 0.1829 - accuracy:
0.9294 - val_loss: 0.4446 - val_accuracy: 0.8482

Epoch 13/75

14/14 [=====] - 0s 4ms/step - loss: 0.1818 - accuracy:
0.9441 - val_loss: 0.4274 - val_accuracy: 0.8482

Epoch 14/75
14/14 [=====] - 0s 17ms/step - loss: 0.2083 - accuracy: 0.9051 - val_loss: 0.4668 - val_accuracy: 0.8482

Epoch 15/75
14/14 [=====] - 0s 4ms/step - loss: 0.2039 - accuracy: 0.9206 - val_loss: 0.4463 - val_accuracy: 0.8661

Epoch 16/75
14/14 [=====] - 0s 4ms/step - loss: 0.2078 - accuracy: 0.9237 - val_loss: 0.4705 - val_accuracy: 0.8482

Epoch 17/75
14/14 [=====] - 0s 4ms/step - loss: 0.2120 - accuracy: 0.9124 - val_loss: 0.4740 - val_accuracy: 0.8304

Epoch 18/75
14/14 [=====] - 0s 4ms/step - loss: 0.1867 - accuracy: 0.9286 - val_loss: 0.4473 - val_accuracy: 0.8393

Epoch 19/75
14/14 [=====] - 0s 4ms/step - loss: 0.2027 - accuracy: 0.9306 - val_loss: 0.4712 - val_accuracy: 0.8393

Epoch 20/75
14/14 [=====] - 0s 4ms/step - loss: 0.1701 - accuracy: 0.9365 - val_loss: 0.4676 - val_accuracy: 0.8393

Epoch 21/75
14/14 [=====] - 0s 4ms/step - loss: 0.2039 - accuracy: 0.8971 - val_loss: 0.4978 - val_accuracy: 0.8482

Epoch 22/75
14/14 [=====] - 0s 4ms/step - loss: 0.1922 - accuracy: 0.9220 - val_loss: 0.4740 - val_accuracy: 0.8214

Epoch 23/75
14/14 [=====] - 0s 4ms/step - loss: 0.1739 - accuracy: 0.9378 - val_loss: 0.5477 - val_accuracy: 0.8304

Epoch 24/75
14/14 [=====] - 0s 4ms/step - loss: 0.1701 - accuracy: 0.9473 - val_loss: 0.5498 - val_accuracy: 0.8661

Epoch 25/75
14/14 [=====] - 0s 4ms/step - loss: 0.1994 - accuracy: 0.9387 - val_loss: 0.5139 - val_accuracy: 0.8304

Epoch 26/75
14/14 [=====] - 0s 5ms/step - loss: 0.1665 - accuracy: 0.9347 - val_loss: 0.5298 - val_accuracy: 0.8304

Epoch 27/75
14/14 [=====] - 0s 4ms/step - loss: 0.1556 - accuracy: 0.9456 - val_loss: 0.5018 - val_accuracy: 0.8304

Epoch 28/75
14/14 [=====] - 0s 5ms/step - loss: 0.1420 - accuracy: 0.9409 - val_loss: 0.5326 - val_accuracy: 0.8214

Epoch 29/75
14/14 [=====] - 0s 5ms/step - loss: 0.1637 - accuracy: 0.9346 - val_loss: 0.5236 - val_accuracy: 0.8393

Epoch 30/75
14/14 [=====] - 0s 4ms/step - loss: 0.1896 - accuracy:
0.9303 - val_loss: 0.4957 - val_accuracy: 0.8482
Epoch 31/75
14/14 [=====] - 0s 4ms/step - loss: 0.1337 - accuracy:
0.9644 - val_loss: 0.5546 - val_accuracy: 0.8393
Epoch 32/75
14/14 [=====] - 0s 4ms/step - loss: 0.1699 - accuracy:
0.9397 - val_loss: 0.5585 - val_accuracy: 0.8214
Epoch 33/75
14/14 [=====] - 0s 4ms/step - loss: 0.1574 - accuracy:
0.9406 - val_loss: 0.5259 - val_accuracy: 0.8393
Epoch 34/75
14/14 [=====] - 0s 4ms/step - loss: 0.1597 - accuracy:
0.9335 - val_loss: 0.5181 - val_accuracy: 0.8393
Epoch 35/75
14/14 [=====] - 0s 5ms/step - loss: 0.1235 - accuracy:
0.9549 - val_loss: 0.5721 - val_accuracy: 0.8482
Epoch 36/75
14/14 [=====] - 0s 4ms/step - loss: 0.1232 - accuracy:
0.9466 - val_loss: 0.5770 - val_accuracy: 0.8393
Epoch 37/75
14/14 [=====] - 0s 4ms/step - loss: 0.1449 - accuracy:
0.9309 - val_loss: 0.5766 - val_accuracy: 0.8482
Epoch 38/75
14/14 [=====] - 0s 4ms/step - loss: 0.1376 - accuracy:
0.9497 - val_loss: 0.5506 - val_accuracy: 0.8571
Epoch 39/75
14/14 [=====] - 0s 4ms/step - loss: 0.1226 - accuracy:
0.9436 - val_loss: 0.5997 - val_accuracy: 0.8482
Epoch 40/75
14/14 [=====] - 0s 4ms/step - loss: 0.1654 - accuracy:
0.9313 - val_loss: 0.5790 - val_accuracy: 0.8482
Epoch 41/75
14/14 [=====] - 0s 4ms/step - loss: 0.1431 - accuracy:
0.9521 - val_loss: 0.6055 - val_accuracy: 0.8482
Epoch 42/75
14/14 [=====] - 0s 4ms/step - loss: 0.1230 - accuracy:
0.9469 - val_loss: 0.5923 - val_accuracy: 0.8482
Epoch 43/75
14/14 [=====] - 0s 4ms/step - loss: 0.1026 - accuracy:
0.9500 - val_loss: 0.6488 - val_accuracy: 0.8393
Epoch 44/75
14/14 [=====] - 0s 15ms/step - loss: 0.1124 - accuracy:
0.9460 - val_loss: 0.6342 - val_accuracy: 0.8482
Epoch 45/75
14/14 [=====] - 0s 4ms/step - loss: 0.1037 - accuracy:
0.9615 - val_loss: 0.6603 - val_accuracy: 0.8304

Epoch 46/75
14/14 [=====] - 0s 4ms/step - loss: 0.1153 - accuracy:
0.9576 - val_loss: 0.6443 - val_accuracy: 0.8393
Epoch 47/75
14/14 [=====] - 0s 4ms/step - loss: 0.0979 - accuracy:
0.9671 - val_loss: 0.6694 - val_accuracy: 0.8571
Epoch 48/75
14/14 [=====] - 0s 4ms/step - loss: 0.0890 - accuracy:
0.9657 - val_loss: 0.6707 - val_accuracy: 0.8393
Epoch 49/75
14/14 [=====] - 0s 4ms/step - loss: 0.0930 - accuracy:
0.9748 - val_loss: 0.6537 - val_accuracy: 0.8304
Epoch 50/75
14/14 [=====] - 0s 5ms/step - loss: 0.1113 - accuracy:
0.9600 - val_loss: 0.6483 - val_accuracy: 0.8482
Epoch 51/75
14/14 [=====] - 0s 4ms/step - loss: 0.0931 - accuracy:
0.9636 - val_loss: 0.6393 - val_accuracy: 0.8482
Epoch 52/75
14/14 [=====] - 0s 4ms/step - loss: 0.0871 - accuracy:
0.9611 - val_loss: 0.6828 - val_accuracy: 0.8482
Epoch 53/75
14/14 [=====] - 0s 4ms/step - loss: 0.1029 - accuracy:
0.9557 - val_loss: 0.7635 - val_accuracy: 0.8393
Epoch 54/75
14/14 [=====] - 0s 4ms/step - loss: 0.0819 - accuracy:
0.9716 - val_loss: 0.8066 - val_accuracy: 0.8304
Epoch 55/75
14/14 [=====] - 0s 4ms/step - loss: 0.0999 - accuracy:
0.9548 - val_loss: 0.7401 - val_accuracy: 0.8571
Epoch 56/75
14/14 [=====] - 0s 4ms/step - loss: 0.0921 - accuracy:
0.9555 - val_loss: 0.7453 - val_accuracy: 0.8393
Epoch 57/75
14/14 [=====] - 0s 5ms/step - loss: 0.1081 - accuracy:
0.9558 - val_loss: 0.7492 - val_accuracy: 0.8482
Epoch 58/75
14/14 [=====] - 0s 4ms/step - loss: 0.0944 - accuracy:
0.9598 - val_loss: 0.7757 - val_accuracy: 0.8482
Epoch 59/75
14/14 [=====] - 0s 4ms/step - loss: 0.1053 - accuracy:
0.9609 - val_loss: 0.8660 - val_accuracy: 0.8393
Epoch 60/75
14/14 [=====] - 0s 4ms/step - loss: 0.1051 - accuracy:
0.9492 - val_loss: 0.8125 - val_accuracy: 0.8482
Epoch 61/75
14/14 [=====] - 0s 4ms/step - loss: 0.0812 - accuracy:
0.9766 - val_loss: 0.7916 - val_accuracy: 0.8482


```

Epoch 62/75
14/14 [=====] - 0s 4ms/step - loss: 0.0870 - accuracy:
0.9601 - val_loss: 0.8041 - val_accuracy: 0.8571
Epoch 63/75
14/14 [=====] - 0s 4ms/step - loss: 0.0835 - accuracy:
0.9647 - val_loss: 0.8099 - val_accuracy: 0.8482
Epoch 64/75
14/14 [=====] - 0s 4ms/step - loss: 0.0963 - accuracy:
0.9639 - val_loss: 0.8542 - val_accuracy: 0.8482
Epoch 65/75
14/14 [=====] - 0s 4ms/step - loss: 0.0626 - accuracy:
0.9793 - val_loss: 0.8070 - val_accuracy: 0.8304
Epoch 66/75
14/14 [=====] - 0s 4ms/step - loss: 0.0732 - accuracy:
0.9657 - val_loss: 0.8366 - val_accuracy: 0.8482
Epoch 67/75
14/14 [=====] - 0s 5ms/step - loss: 0.0623 - accuracy:
0.9800 - val_loss: 0.8704 - val_accuracy: 0.8482
Epoch 68/75
14/14 [=====] - 0s 4ms/step - loss: 0.1131 - accuracy:
0.9484 - val_loss: 0.8432 - val_accuracy: 0.8482
Epoch 69/75
14/14 [=====] - 0s 4ms/step - loss: 0.0841 - accuracy:
0.9647 - val_loss: 0.8692 - val_accuracy: 0.8482
Epoch 70/75
14/14 [=====] - 0s 4ms/step - loss: 0.0869 - accuracy:
0.9696 - val_loss: 0.8658 - val_accuracy: 0.8393
Epoch 71/75
14/14 [=====] - 0s 4ms/step - loss: 0.0691 - accuracy:
0.9737 - val_loss: 0.8994 - val_accuracy: 0.8304
Epoch 72/75
14/14 [=====] - 0s 4ms/step - loss: 0.0608 - accuracy:
0.9781 - val_loss: 0.9170 - val_accuracy: 0.8482
Epoch 73/75
14/14 [=====] - 0s 4ms/step - loss: 0.0627 - accuracy:
0.9796 - val_loss: 0.8816 - val_accuracy: 0.8482
Epoch 74/75
14/14 [=====] - 0s 4ms/step - loss: 0.0850 - accuracy:
0.9749 - val_loss: 0.9836 - val_accuracy: 0.8304
Epoch 75/75
14/14 [=====] - 0s 4ms/step - loss: 0.0940 - accuracy:
0.9472 - val_loss: 0.9276 - val_accuracy: 0.8304

```

```

[ ]: hist = pd.DataFrame(history.history)
hist['epoch'] = history.epoch
hist.tail()

```

```
[ ]:      loss  accuracy  val_loss  val_accuracy  epoch
70  0.083450  0.963964  0.899368      0.830357    70
71  0.080079  0.970721  0.917036      0.848214    71
72  0.079431  0.968468  0.881565      0.848214    72
73  0.072801  0.970721  0.983648      0.830357    73
74  0.081346  0.961712  0.927632      0.830357    74
```

```
[ ]: acc = (hist['accuracy'].tail().sum())*100/5
      val_acc = (hist['val_accuracy'].tail().sum())*100/5

      print("Training Accuracy = {}% and Validation Accuracy= {}%".
            ↪format(acc,val_acc))
```

Training Accuracy = 96.71171188354492% and Validation Accuracy= 83.7499988079071%

```
[ ]: y_pred = model.predict_classes(X_test)
      print("----- Accuracy ----- \n")
      print(accuracy_score(y_test,y_pred))
      print('-----')
      print("----- Confusion Matrix ----- \n")
      print(confusion_matrix(y_test,y_pred))
      print("----- Classification Report----- \n")
      print(classification_report(y_test,y_pred))
      print('-----')
```

----- Accuracy -----

0.7571428571428571

----- Confusion Matrix -----

```
[[70 22]
 [12 36]]
```

----- Classification Report-----

	precision	recall	f1-score	support
0	0.85	0.76	0.80	92
1	0.62	0.75	0.68	48
accuracy			0.76	140
macro avg	0.74	0.76	0.74	140
weighted avg	0.77	0.76	0.76	140

/usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/engine/sequential.py:450: UserWarning:

`model.predict_classes()` is deprecated and will be removed after 2021-01-01. Please use instead: `* np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). `*(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).
`warnings.warn('model.predict_classes() is deprecated and '`

14 Conclusion

Best accuracy was by using **XGboost Classifier** with 93.4% accuracy, and an excellent validation accuracy (89.5 %). The accuracy achieved by using: 1. Forward filling the NaN, 2. without using the extracted feature, 3. DBSCAN Density-Based Spatial Clustering of Applications for outlier detection, 4. Random over sampler to balance the classes

15 Step 6: Reproducibility

```
[!]: !pip install watermark
```

Collecting watermark

Downloading <https://files.pythonhosted.org/packages/60/fe/3ed83b6122e70dce6fe269dfd763103c333f168bf91037add73ea4fe81c2/watermark-2.0.2-py2.py3-none-any.whl>
Requirement already satisfied: ipython in /usr/local/lib/python3.6/dist-packages (from watermark) (5.5.0)

Requirement already satisfied: simplegeneric>0.8 in /usr/local/lib/python3.6/dist-packages (from ipython->watermark) (0.8.1)

Requirement already satisfied: decorator in /usr/local/lib/python3.6/dist-packages (from ipython->watermark) (4.4.2)

Requirement already satisfied: pygments in /usr/local/lib/python3.6/dist-packages (from ipython->watermark) (2.6.1)

Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.6/dist-packages (from ipython->watermark) (4.3.3)

Requirement already satisfied: pexpect; sys_platform != "win32" in /usr/local/lib/python3.6/dist-packages (from ipython->watermark) (4.8.0)

Requirement already satisfied: prompt-toolkit<2.0.0,>=1.0.4 in /usr/local/lib/python3.6/dist-packages (from ipython->watermark) (1.0.18)

Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.6/dist-packages (from ipython->watermark) (51.0.0)

Requirement already satisfied: pickleshare in /usr/local/lib/python3.6/dist-packages (from ipython->watermark) (0.7.5)

Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from traitlets>=4.2->ipython->watermark) (1.15.0)

Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.6/dist-packages (from traitlets>=4.2->ipython->watermark) (0.2.0)

Requirement already satisfied: ptyprocess>=0.5 in /usr/local/lib/python3.6/dist-packages (from pexpect; sys_platform != "win32"->ipython->watermark) (0.6.0)

```
Requirement already satisfied: wcwidth in /usr/local/lib/python3.6/dist-packages
(from prompt-toolkit<2.0.0,>=1.0.4->ipython->watermark) (0.2.5)
Installing collected packages: watermark
Successfully installed watermark-2.0.2
```

```
[ ]: %load_ext watermark
```

```
[ ]: %watermark -a 'Hazim' -nmv --packages numpy,pandas,sklearn,matplotlib,seaborn
```

Hazim Tue Dec 29 2020

CPython 3.6.9

IPython 5.5.0

numpy 1.19.4

pandas 1.1.5

sklearn 0.0

matplotlib 3.2.2

seaborn 0.11.0

compiler : GCC 8.4.0

system : Linux

release : 4.19.112+

machine : x86_64

processor : x86_64

CPU cores : 2

interpreter: 64bit